

EfficientML.ai Lecture 06

Quantization

Part II



Song Han

Associate Professor, MIT
Distinguished Scientist, NVIDIA

 @SongHan_MIT



Lecture Plan

Today we will:

1. Review Linear Quantization.
2. Introduce **Post-Training Quantization (PTQ)** that quantizes a floating-point neural network model, including: channel quantization, group quantization, and range clipping.
3. Introduce **Quantization-Aware Training (QAT)** that emulates inference-time quantization during the training/fine-tuning and recover the accuracy.
4. Introduce **binary and ternary** quantization.
5. Introduce automatic **mixed-precision** quantization.

Neural Network Quantization

2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

3	0	2	1	3:	2.00
1	1	0	3	2:	1.50
0	3	1	0	1:	0.00
3	1	2	2	0:	-1.00

1	-2	0	-1
-1	-1	-2	1
-2	1	-1	-2
1	-1	0	0


(- -1) × 1.07

K-Means-based
Quantization

Linear
Quantization

Storage	Floating-Point Weights	Integer Weights; Floating-Point Codebook	Integer Weights
Computation	Floating-Point Arithmetic	Floating-Point Arithmetic	Integer Arithmetic


K-Means-based Weight Quantization



Deep Compression [Han et al., ICLR 2016]

K-Means-based Weight Quantization

Accuracy vs. compression rate for AlexNet on ImageNet dataset



Deep Compression [Han *et al.*, ICLR 2016]

Linear Quantization

An affine mapping of integers to real numbers $r = S(q - Z)$

weights
(32-bit float)

2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

quantized weights
(2-bit signed int)

1	-2	0	-1
-1	-1	-2	1
-2	1	-1	-2
1	-1	0	0

zero point
(2-bit signed int)

-1

scale
(32-bit float)


1.07

2.14	-1.07	1.07	0
0	0	-1.07	2.14
-1.07	2.14	0	-1.07
2.14	0	1.07	1.07

Binary	Decimal
01	1
00	0
11	-1
10	-2

Linear Quantization

An affine mapping of integers to real numbers $r = S(q - Z)$



Bit Width	q_{\min}	q_{\max}
2	-2	1
3	-4	3
4	-8	7
N	-2^{N-1}	$2^{N-1}-1$

Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference [Jacob *et al.*, CVPR 2018]

Linear Quantized Fully-Connected Layer

Linear Quantization is an affine mapping of integers to real numbers $r = S(q - Z)$

- Consider the following fully-connected layer.

$$\mathbf{Y} = \mathbf{W}\mathbf{X} + \mathbf{b}$$

$$Z_W = 0$$

$$Z_b = 0, \quad S_b = S_W S_X$$

$$\mathbf{q}_{bias} = \mathbf{q}_b - Z_X \mathbf{q}_W$$

$$\mathbf{q}_Y = \boxed{\frac{S_W S_X}{S_Y}} \left(\boxed{\mathbf{q}_W \mathbf{q}_X} + \boxed{\mathbf{q}_{bias}} \right) + \boxed{Z_Y}$$

Rescale to N -bit Int N -bit Int Mult.
 N -bit Int 32-bit Int Add. N -bit Int Add

Note: both \mathbf{q}_b and \mathbf{q}_{bias} are 32 bits.

Linear Quantized Convolution Layer

Linear Quantization is an affine mapping of integers to real numbers $r = S(q - Z)$

- Consider the following convolution layer.

$$Y = \text{Conv}(W, X) + b$$

$$Z_W = 0$$

$$Z_b = 0, \quad S_b = S_W S_X$$

$$q_{bias} = q_b - \text{Conv}(q_W, Z_X)$$


$$q_Y = \frac{S_W S_X}{S_Y} \left(\text{Conv}(q_W, q_X) + q_{bias} \right) + Z_Y$$

Rescale to N -bit Int N -bit Int Mult. 32-bit Int Add. N -bit Int Add

Note: both q_b and q_{bias} are 32 bits.

Scale and Zero Point of Linear Quantization

Linear Quantization is an affine mapping of integers to real numbers $r = S(q - Z)$




2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

$$\begin{aligned} S &= \frac{r_{\max} - r_{\min}}{q_{\max} - q_{\min}} \\ &= \frac{2.12 - (-1.08)}{1 - (-2)} \\ &= 1.07 \end{aligned}$$

$$\begin{aligned} Z &= q_{\min} - \frac{r_{\min}}{S} \\ &= \text{round}\left(-2 - \frac{-1.08}{1.07}\right) \\ &= -1 \end{aligned}$$

Scale and Zero Point of Linear Quantization

Linear Quantization is an affine mapping of integers to real numbers $r = S(q - Z)$



2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

$$\begin{aligned} S &= \frac{|r|_{\max}}{q_{\max}} \\ &= \frac{2.12}{1} \\ &= 2.12 \end{aligned}$$

$$Z = 0$$

Post-Training Quantization

How should we get the optimal linear quantization parameters (S, Z)?

Topic I: Quantization Granularity

Topic II: Dynamic Range Clipping

Topic III: Rounding

Post-Training Quantization

How should we get the optimal linear quantization parameters (S, Z)?

Topic I: Quantization Granularity

Topic II: Dynamic Range Clipping


Topic III: Rounding

Quantization Granularity


- Per-Tensor Quantization



- Per-Channel Quantization



- Group Quantization



- Per-Vector Quantization


- Shared Micro-exponent (MX) data type

Quantization Granularity


- **Per-Tensor Quantization**



- Per-Channel Quantization





- Group Quantization



- Per-Vector Quantization

- Shared Micro-exponent (MX) data type

Symmetric Linear Quantization on Weights



- $|r|_{\max} = |\mathbf{W}|_{\max}$
- Using *single* scale S for whole weight tensor (**Per-Tensor Quantization**)
 - works well for large models
 - accuracy drops for small models
- Common failure results from
 - large differences (more than 100×) in ranges of weights for different output channels — outlier weight
- Solution: **Per-Channel Quantization**

Data-Free Quantization Through Weight Equalization and Bias Correction [Markus *et al.*, ICCV 2019]


Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference [Jacob *et al.*, CVPR 2018]

Quantization Granularity


- Per-Tensor Quantization



- **Per-Channel Quantization**



- Group Quantization



- Per-Vector Quantization

- Shared Micro-exponent (MX) data type

Per-Channel Weight Quantization

Example: 2-bit linear quantization

		Per-Channel Quantization				Per-Tensor Quantization			
		iC							
OC		2.09	-0.98	1.48	0.09				
		0.05	-0.14	-1.08	2.12				
		-0.91	1.92	0	-1.03				
		1.87	0	1.53	1.49				

Per-Channel Weight Quantization

Example: 2-bit linear quantization

		Per-Channel Quantization			
		ic			
OC		2.09	-0.98	1.48	0.09
		0.05	-0.14	-1.08	2.12
		-0.91	1.92	0	-1.03
		1.87	0	1.53	1.49

Per-Tensor Quantization

$$|r|_{\max} = 2.12$$

$$S = \frac{|r|_{\max}}{q_{\max}} = \frac{2.12}{2^{2-1} - 1} = 2.12$$

1	0	1	0	2.12	0	2.12	0
0	0	-1	1	0	0	-2.12	2.12
0	1	0	0	0	2.12	0	0
1	0	1	1	2.12	0	2.12	2.12

Quantized

Reconstructed

$$\|\mathbf{W} - S\mathbf{q}_W\|_F = 2.28$$

Per-Channel Weight Quantization

Example: 2-bit linear quantization

		Per-Channel Quantization			
		ic			
OC		2.09	-0.98	1.48	0.09
		0.05	-0.14	-1.08	2.12
		-0.91	1.92	0	-1.03
		1.87	0	1.53	1.49
		$ r _{\max} = 2.09$	$ r _{\max} = 2.12$	$ r _{\max} = 1.92$	$ r _{\max} = 1.87$
		$S_0 = 2.09$	$S_1 = 2.12$	$S_2 = 1.92$	$S_3 = 1.87$

Per-Tensor Quantization

$$|r|_{\max} = 2.12$$

$$S = \frac{|r|_{\max}}{q_{\max}} = \frac{2.12}{2^{2-1} - 1} = 2.12$$

1	0	1	0
0	0	-1	1
0	1	0	0
1	0	1	1

2.12	0	2.12	0
0	0	-2.12	2.12
0	2.12	0	0
2.12	0	2.12	2.12

Quantized

Reconstructed

$$\|\mathbf{W} - S\mathbf{q}_W\|_F = 2.28$$

Per-Channel Weight Quantization

Example: 2-bit linear quantization

		Per-Channel Quantization			
		ic			
OC		2.09	-0.98	1.48	0.09
		0.05	-0.14	-1.08	2.12
		-0.91	1.92	0	-1.03
		1.87	0	1.53	1.49
		$ r _{\max} = 2.09$		$S_0 = 2.09$	
		$ r _{\max} = 2.12$		$S_1 = 2.12$	
		$ r _{\max} = 1.92$		$S_2 = 1.92$	
		$ r _{\max} = 1.87$		$S_3 = 1.87$	

1	0	1	0
0	0	-1	1
0	1	0	-1
1	0	1	1

Quantized

2.09	0	2.09	0
0	0	-2.12	2.12
0	1.92	0	-1.92
1.87	0	1.87	1.87

Reconstructed

$$\|\mathbf{W} - \mathbf{S} \odot \mathbf{q}_W\|_F = 2.08$$

Per-Tensor Quantization

$$|r|_{\max} = 2.12$$

$$S = \frac{|r|_{\max}}{q_{\max}} = \frac{2.12}{2^{2-1} - 1} = 2.12$$

1	0	1	0
0	0	-1	1
0	1	0	0
1	0	1	1

Quantized

2.12	0	2.12	0
0	0	-2.12	2.12
0	2.12	0	0
2.12	0	2.12	2.12

Reconstructed

$$\|\mathbf{W} - \mathbf{S} \mathbf{q}_W\|_F = 2.28$$

Per-Channel Weight Quantization

Example: 2-bit linear quantization

		Per-Channel Quantization			
		ic			
OC		2.09	-0.98	1.48	0.09
		0.05	-0.14	-1.08	2.12
		-0.91	1.92	0	-1.03
		1.87	0	1.53	1.49
		$ r _{\max} = 2.09$		$S_0 = 2.09$	
		$ r _{\max} = 2.12$		$S_1 = 2.12$	
		$ r _{\max} = 1.92$		$S_2 = 1.92$	
		$ r _{\max} = 1.87$		$S_3 = 1.87$	

1	0	1	0
0	0	-1	1
0	1	0	-1
1	0	1	1

Quantized

2.09	0	2.09	0
0	0	-2.12	2.12
0	1.92	0	-1.92
1.87	0	1.87	1.87

Reconstructed

$$\|\mathbf{W} - \mathbf{S} \odot \mathbf{q}_W\|_F = 2.08$$

Per-Tensor Quantization

$$|r|_{\max} = 2.12$$
$$S = \frac{|r|_{\max}}{q_{\max}} = \frac{2.12}{2^{2-1} - 1} = 2.12$$

1	0	1	0
0	0	-1	1
0	1	0	0
1	0	1	1

Quantized

2.12	0	2.12	0
0	0	-2.12	2.12
0	2.12	0	0
2.12	0	2.12	2.12

Reconstructed


$$\|\mathbf{W} - \mathbf{S} \mathbf{q}_W\|_F = 2.28$$

Quantization Granularity


- Per-Tensor Quantization



- Per-Channel Quantization



- **Group Quantization**




- **Per-Vector Quantization**

- **Shared Micro-exponent (MX) data type**

VS-Quant: Per-vector Scaled Quantization

Hierarchical scaling factor

- $r = S(q - Z) \rightarrow r = \gamma \cdot S_q(q - Z)$
 - γ is a floating-point coarse grained scale factor
 - S_q is an integer per-vector scale factor
 - achieves a balance between accuracy and hardware efficiency by
 - less expensive integer scale factors at finer granularity
 - more expensive floating-point scale factors at coarser granularity
- Memory Overhead of two-level scaling:
 - Given 4-bit quantization with 4-bit per-vector scale for every 16 elements, the effective bit width is $4 + 4 / 16 = 4.25$ bits.



Group Quantization

Multi-level scaling scheme

	w_{11}						

$$r = (q - z) \cdot s \rightarrow$$

$$r = (q - z) \cdot s_{l_0} \cdot s_{l_1} \cdot \dots$$

r : real number value


q : quantized value

z : zero point ($z = 0$ is symmetric quantization)

s : scale factors of different levels


Group Quantization

Multi-level scaling scheme



$$r = (q - z) \cdot s_{l_0} \cdot s_{l_1} \cdot \dots$$


- r : real number value
- q : quantized value
- z : zero point ($z = 0$ is symmetric quantization)
- s : scale factors of different levels



Quantization Approach	Data Type	L0 Group Size	L0 Scale Data Type	L1 Group Size	L1 Scale Data Type	Effective Bit Width
Per-Channel Quant	INT4	Per Channel	FP16	-	-	4

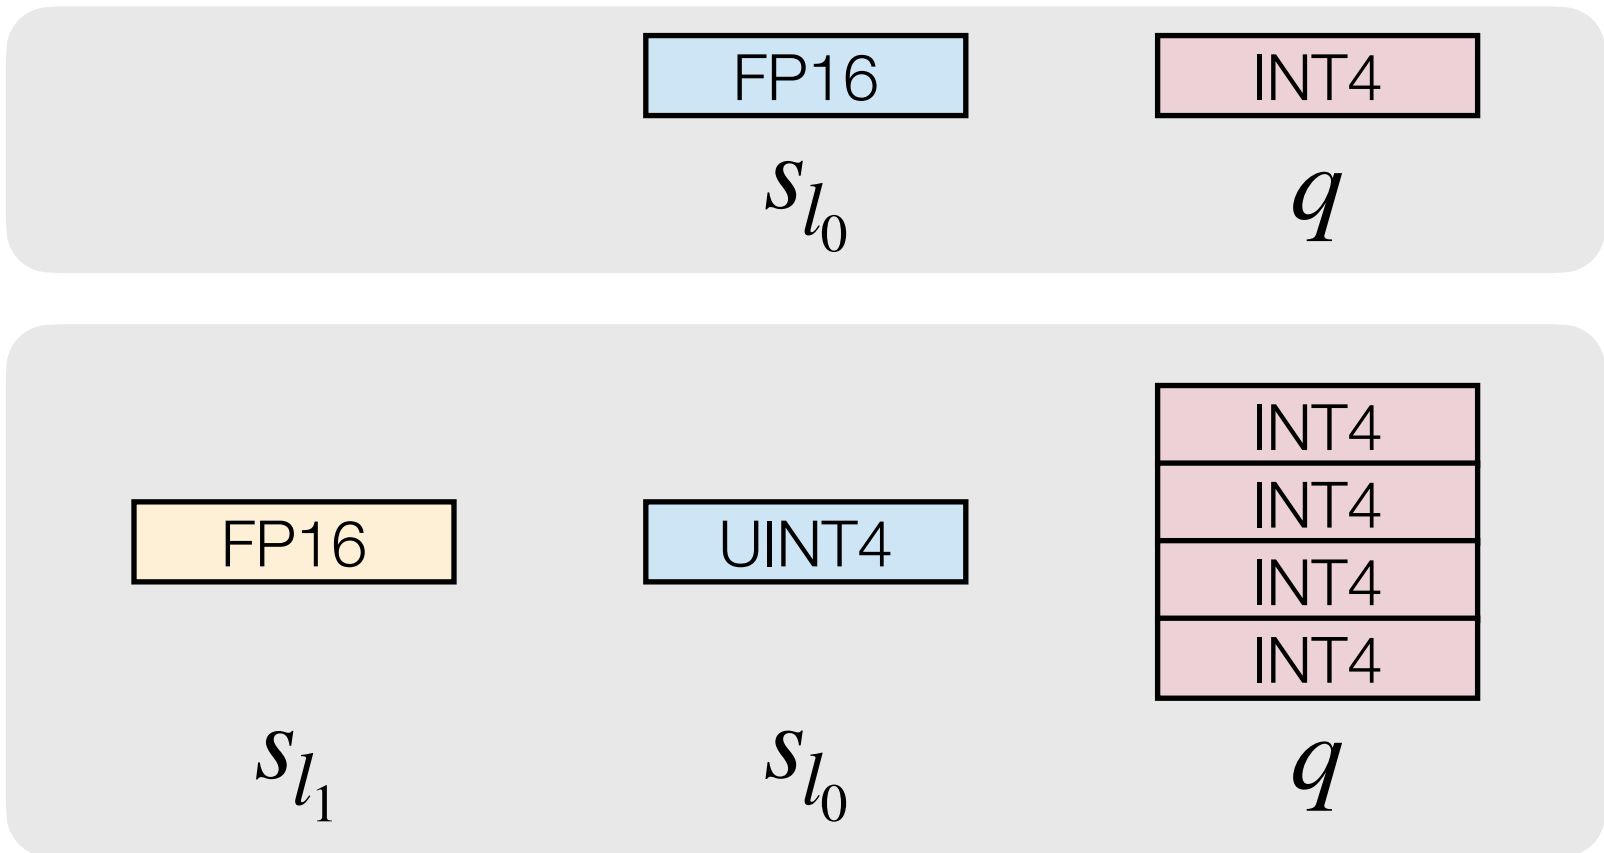
Group Quantization

Multi-level scaling scheme



$$r = (q - z) \cdot s_{l_0} \cdot s_{l_1} \cdot \dots$$


r : real number value
 q : quantized value
 z : zero point ($z = 0$ is symmetric quantization)
 s : scale factors of different levels



Quantization Approach	Data Type	L0 Group Size	L0 Scale Data Type	L1 Group Size	L1 Scale Data Type	Effective Bit Width
Per-Channel Quant	INT4	Per Channel	FP16	-	-	4
VSQ	INT4	16	UINT4	Per Channel	FP16	4+4/16=4.25


Group Quantization

Multi-level scaling scheme



$$r = (q - z) \cdot s_{l_0} \cdot s_{l_1} \cdot \dots$$

r : real number value
 q : quantized value
 z : zero point ($z = 0$ is symmetric quantization)
 s : scale factors of different levels



Quantization Approach	Data Type	L0 Group Size	L0 Scale Data Type	L1 Group Size	L1 Scale Data Type	Effective Bit Width
Per-Channel Quant	INT4	Per Channel	FP16	-	-	4
VSQ	INT4	16	UINT4	Per Channel	FP16	4+4/16=4.25
MX4	S1M2	2	E1M0	16	E8M0	3+1/2+8/16=4
MX6	S1M4	2	E1M0	16	E8M0	5+1/2+8/16=6
MX9	S1M7	2	E1M0	16	E8M0	8+1/2+8/16=9

With Shared Microexponents, A Little Shifting Goes a Long Way [Bita Rouhani et al.]

Post-Training Quantization


How should we get the optimal linear quantization parameters (S, Z)?

Topic I: Quantization Granularity


Topic II: Dynamic Range Clipping

Topic III: Rounding

Linear Quantization on Activations



- Unlike weights, the activation range varies across inputs.
- To determine the floating-point range, the activations statistics are gathered **before** deploying the model.




Dynamic Range for Activation Quantization

Collect activations statistics before deploying the model

$$\hat{r}_{\max, \min}^{(t)} = \alpha \cdot r_{\max, \min}^{(t)} + (1 - \alpha) \cdot \hat{r}_{\max, \min}^{(t-1)}$$



- Type 1: During training
 - Exponential moving averages (EMA)
 - observed ranges are smoothed across thousands of training steps



Dynamic Range for Activation Quantization

Collect activations statistics before deploying the model

- Type 2: By running a few “calibration” batches of samples on the trained FP32 model
- spending dynamic range on the outliers hurts the representation ability.
- use *mean* of the min/max of each sample in the batches
- analytical calculation (see next slide)



Neural Network Distiller

Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference [Jacob *et al.*, CVPR 2018]

Dynamic Range for Activation Quantization

Collect activations statistics before deploying the model

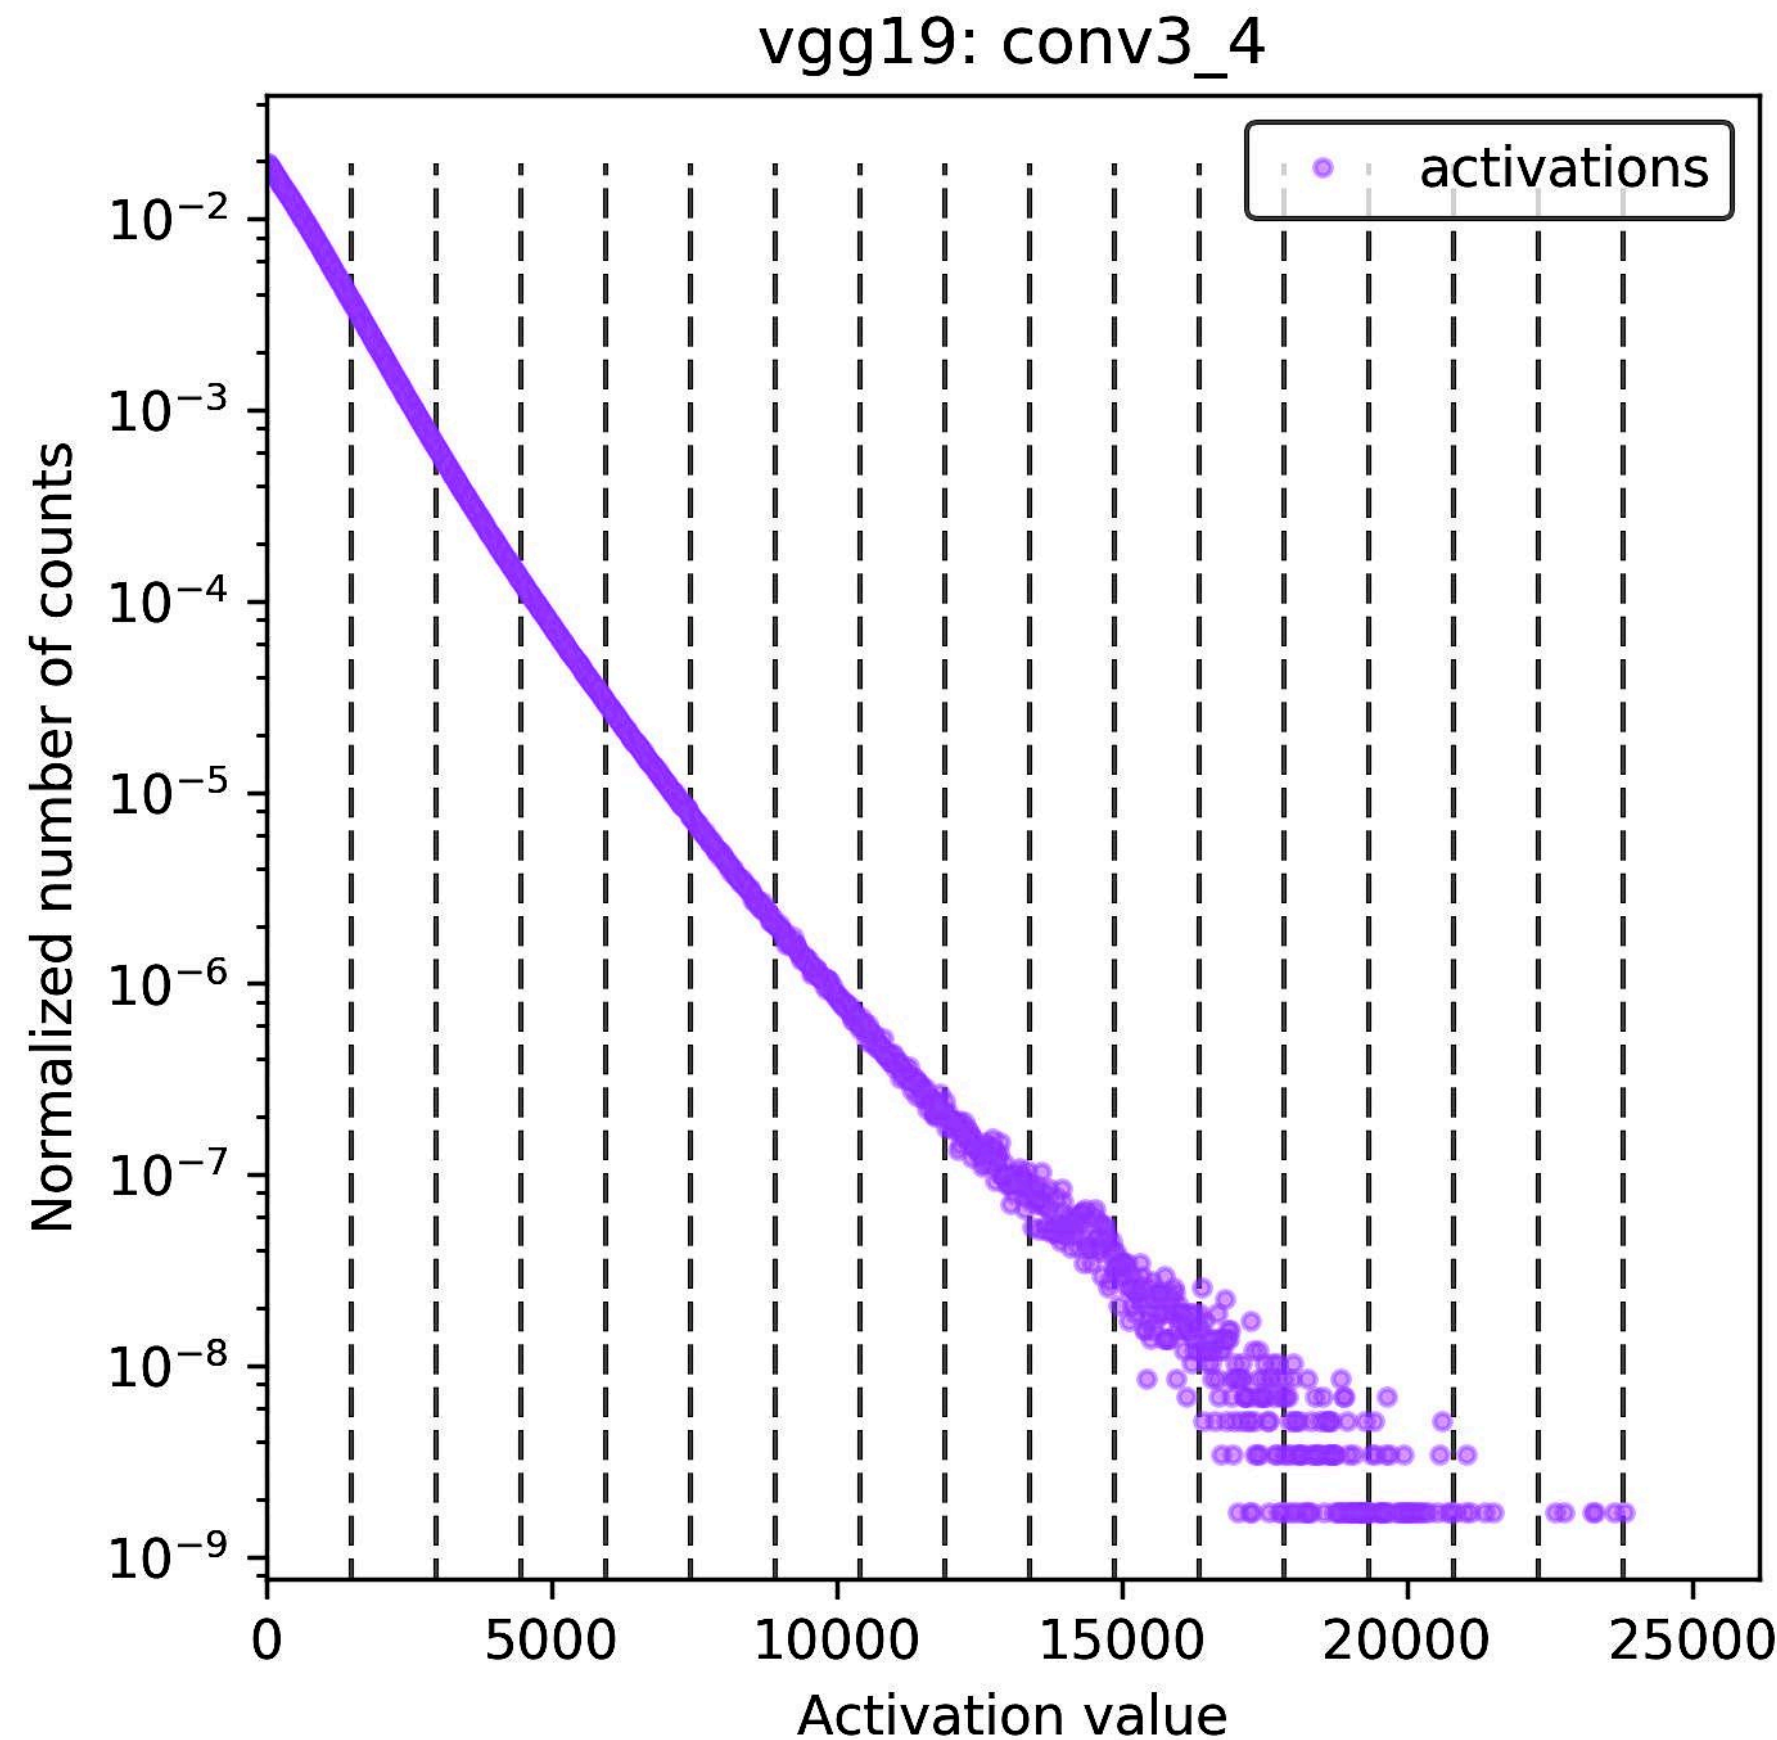
- Type 2: By running a few “calibration” batches of samples on the trained FP32 model
- minimize loss of information, since integer model encodes the same information as the original floating-point model.

- loss of information is measured by **Kullback-Leibler divergence** (relative entropy or information divergence):

- for two discrete probability distributions P, Q

$$D_{KL}(P||Q) = \sum_i^N P(x_i) \log \frac{P(x_i)}{Q(x_i)}$$




- intuition: KL divergence measures the amount of information lost when approximating a given encoding.



8-bit Inference with TensorRT [Szymon Migacz, 2017]

Dynamic Range for Activation Quantization

Minimize loss of information by minimizing the KL divergence

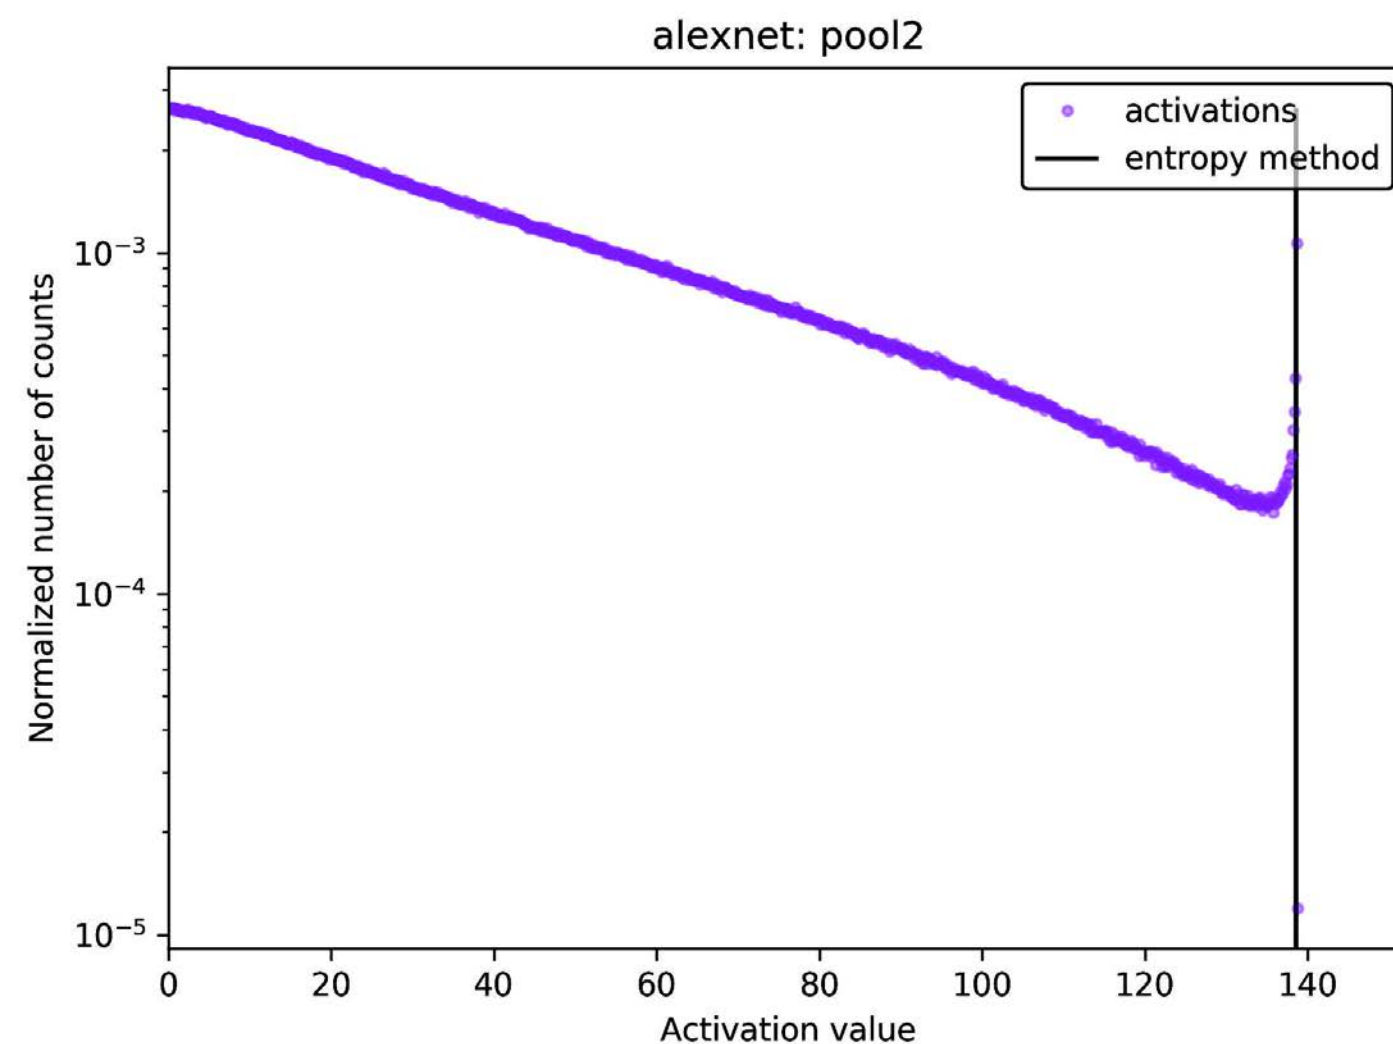


8-bit Inference with TensorRT [Szymon Migacz, 2017]


Dynamic Range for Activation Quantization

Minimize loss of information by minimizing the KL divergence


AlexNet: Pool 2




GoogleNet: incpetion_5a/5x5



ResNet-152: res4b8_branch2a



GoogleNet: incpetion_3a/pool




8-bit Inference with TensorRT [Szymon Migacz, 2017]


Dynamic Range for Quantization

Minimize mean-square-error (MSE) using Newton-Raphson method

max-scaled quantization





clipped quantization



Optimal Clipping and Magnitude-aware Differentiation for Improved Quantization-aware Training [Sakr *et al.*, ICML 2022]

Dynamic Range for Quantization

Minimize mean-square-error (MSE) using Newton-Raphson method



Network	FP32 Accuracy	OCTAV int4
ResNet-50	76.07	75.84
MobileNet-V2	71.71	70.88
Bert-Large	91.00	87.09

Optimal Clipping and Magnitude-aware Differentiation for Improved Quantization-aware Training [Sakr *et al.*, ICML 2022]

Post-Training Quantization

How should we get the optimal linear quantization parameters (S, Z)?

Topic I: Quantization Granularity

Topic II: Dynamic Range Clipping


Topic III: Rounding

Adaptive Rounding for Weight Quantization

Rounding-to-nearest is not optimal

- **Philosophy**

- Rounding-to-nearest is not optimal
- Weights are correlated with each other. The best rounding for each weight (to nearest) is not the best rounding for the whole tensor



- What is optimal? Rounding that reconstructs the original activation the best, which may be very different
 - For weight quantization only
 - With short-term tuning, (almost) post-training quantization

Adaptive Rounding for Weight Quantization

Rounding-to-nearest is not optimal

- **Method:**

- Instead of $\lfloor w \rfloor$, we want to choose from $\{ \lfloor w \rfloor, \lceil w \rceil \}$ to get the best reconstruction
- We took a learning-based method to find quantized value $\tilde{w} = \lfloor \lfloor w \rfloor + \delta \rfloor, \delta \in [0, 1]$

Adaptive Rounding for Weight Quantization


Rounding-to-nearest is not optimal

- **Method:**

- Instead of $\lfloor w \rfloor$, we want to choose from $\{ \lfloor w \rfloor, \lceil w \rceil \}$ to get the best reconstruction
- We took a learning-based method to find quantized value $\tilde{w} = \lfloor \lfloor w \rfloor + \delta \rfloor, \delta \in [0,1]$
- We optimize the following equation (omit the derivation):

$$\begin{aligned} & \operatorname{argmin}_{\mathbf{V}} \|\mathbf{W}\mathbf{x} - \tilde{\mathbf{W}}\mathbf{x}\|_F^2 + \lambda f_{reg}(\mathbf{V}) \\ \rightarrow & \operatorname{argmin}_{\mathbf{V}} \|\mathbf{W}\mathbf{x} - \lfloor \lfloor \mathbf{W} \rfloor + \mathbf{h}(\mathbf{V}) \rfloor \mathbf{x} \|_F^2 + \lambda f_{reg}(\mathbf{V}) \end{aligned}$$

- \mathbf{x} is the input to the layer, \mathbf{V} is a random variable of the same shape
- $\mathbf{h}()$ is a function to map the range to $(0,1)$, such as rectified sigmoid
- $f_{reg}(\mathbf{V})$ is a regularization that encourages $\mathbf{h}(\mathbf{V})$ to be binary



$$f_{reg}(\mathbf{V}) = \sum_{i,j} 1 - |2h(\mathbf{V}_{i,j}) - 1|^\beta$$

Neural Network Quantization

2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

3	0	2	1	3:	2.00
1	1	0	3	2:	1.50
0	3	1	0	1:	0.00
3	1	2	2	0:	-1.00

1	-2	0	-1
-1	-1	-2	1
-2	1	-1	-2
1	-1	0	0

$(\text{matrix}) - (-1) \times 1.07$

K-Means-based
Quantization

Linear
Quantization

Storage	Floating-Point Weights	Integer Weights; Floating-Point Codebook	Integer Weights
Computation	Floating-Point Arithmetic	Floating-Point Arithmetic	Integer Arithmetic

- **Zero Point**
 - Asymmetric
 - Symmetric
- **Scaling Granularity**
 - Per-Tensor
 - Per-Channel
 - Group Quantization
- **Range Clipping**
 - Exponential Moving Average
 - Minimizing KL Divergence
 - Minimizing Mean-Square-Error
- **Rounding**
 - Round-to-Nearest
 - AdaRound

Post-Training INT8 Linear Quantization

Activation		Symmetric	Asymmetric
		Per-Tensor	Per-Tensor
		Minimize KL-Divergence	Exponential Moving Average (EMA)
Weight		Symmetric	Symmetric
		Per-Tensor	Per-Channel
Neural Network	GoogLeNet	-0.45%	0%
	ResNet-50	-0.13%	-0.6%
	ResNet-152	-0.08%	-1.8%
	MobileNetV1	-	-11.8%
	MobileNetV2	-	-2.1%

Data-Free Quantization Through Weight Equalization and Bias Correction [Markus *et al.*, ICCV 2019]
Quantizing Deep Convolutional Networks for Efficient Inference: A Whitepaper [Raghuraman Krishnamoorthi, arXiv 2018]
8-bit Inference with TensorRT [Szymon Migacz, 2017]

Post-Training INT8 Linear Quantization

Activation		Symmetric	Asymmertric
		Per-Tensor	Per-Tensor
		Minimize KL-Divergence	Exponential Moving Average (EMA)
Weight		Symmetric	Symmetric
		Per-Tensor	Per-Channel
Neural Network	Smaller models seem to not respond as well to post-training quantization, presumably due to their smaller representational capacity.		How should we improve performance of quantized models?
	MobileNetV1	-	-11.8%
	MobileNetV2	-	-2.1%

Data-Free Quantization Through Weight Equalization and Bias Correction [Markus *et al.*, ICCV 2019]
Quantizing Deep Convolutional Networks for Efficient Inference: A Whitepaper [Raghuraman Krishnamoorthi, arXiv 2018]
8-bit Inference with TensorRT [Szymon Migacz, 2017]

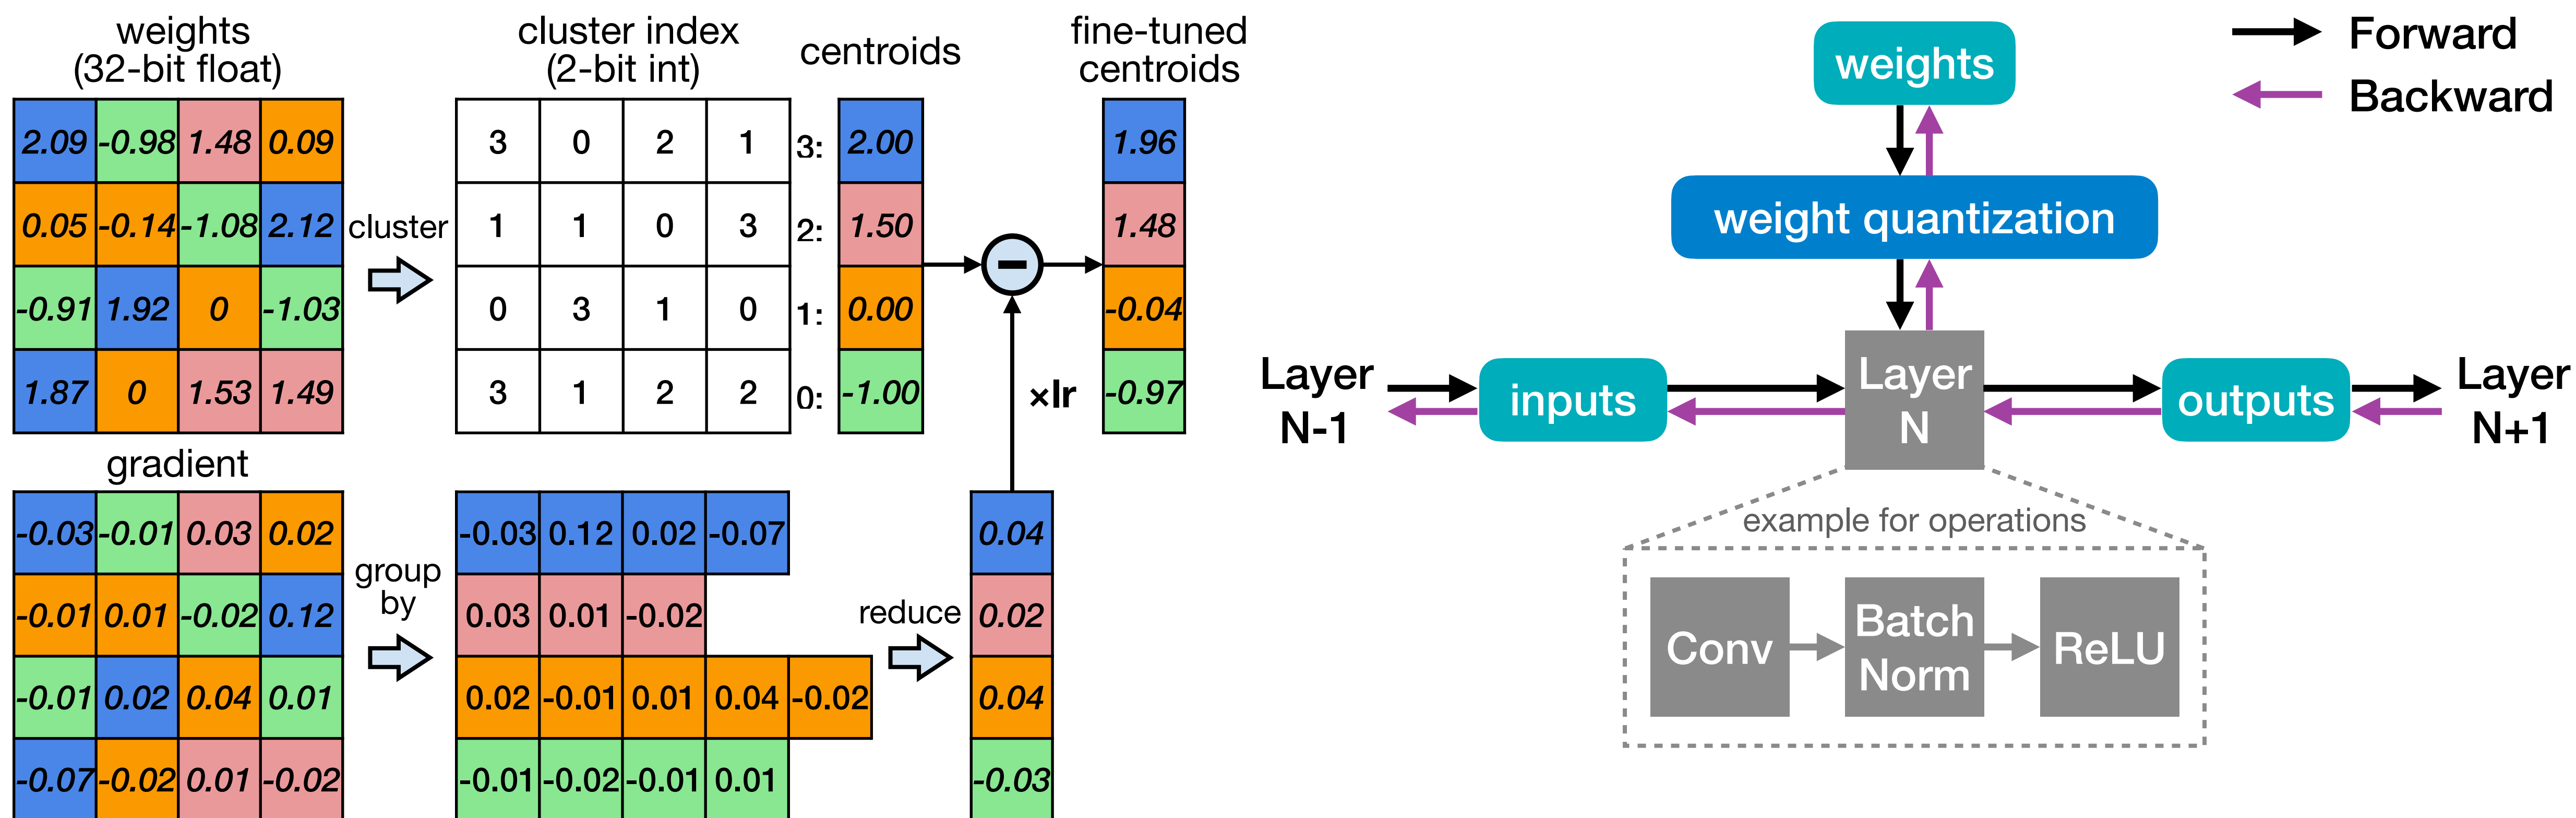
Quantization-Aware Training

How should we improve performance of quantized models?

Quantization-Aware Training

Train the model taking quantization into consideration

- To minimize the loss of accuracy, especially aggressive quantization with 4 bits and lower bit width, neural network will be trained/fine-tuned with quantized weights and activations.
- Usually, fine-tuning a pre-trained floating point model provides better accuracy than training from scratch.




Deep Compression [Han et al., ICLR 2016]

Quantization-Aware Training

Train the model taking quantization into consideration

- A full precision copy of the weights W is maintained throughout the training.
- The small gradients are accumulated without loss of precision.
- Once the model is trained, only the quantized weights are used for inference.

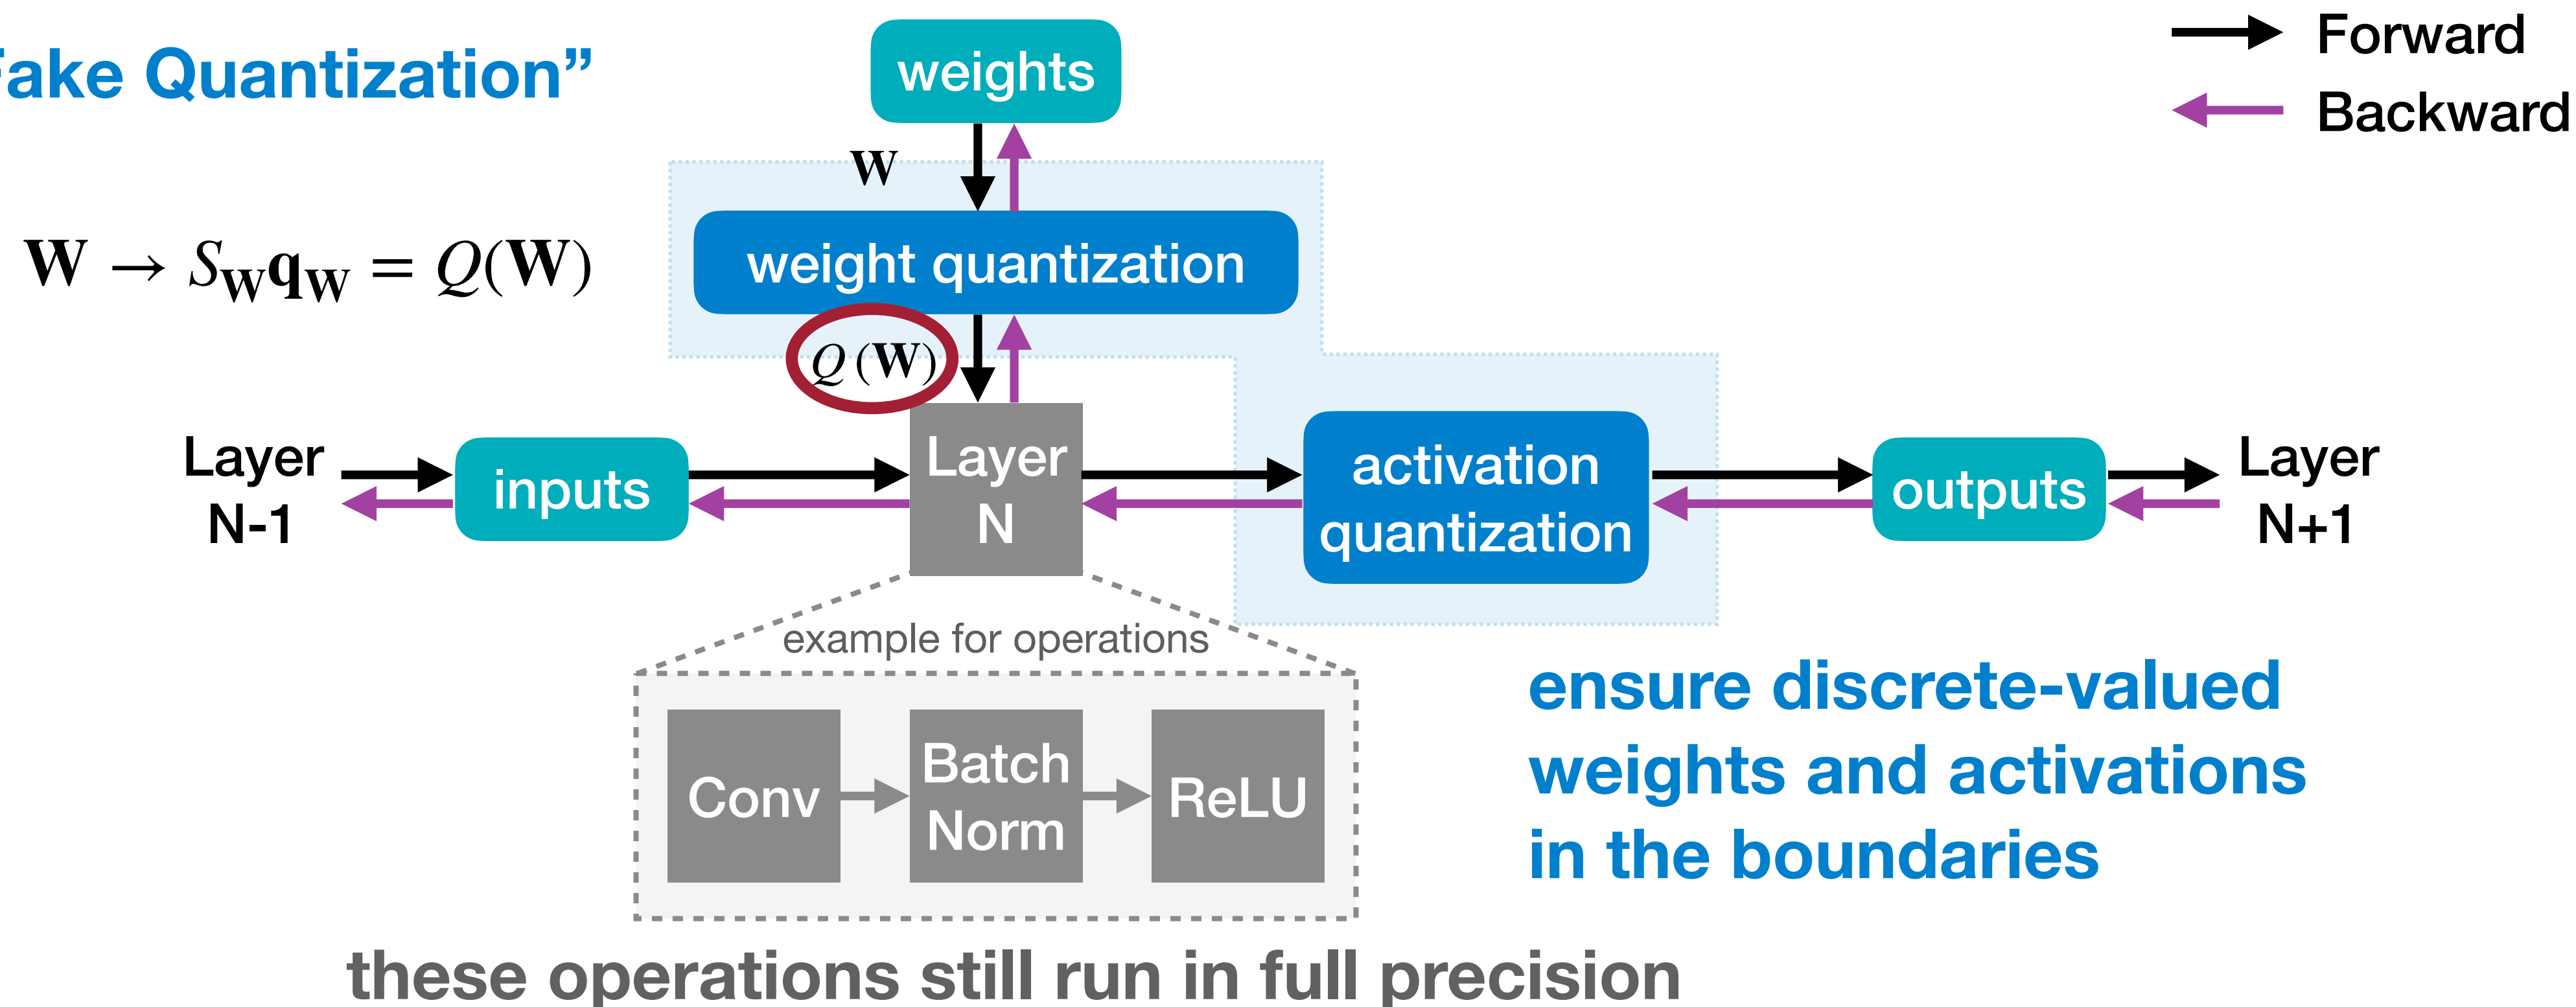


Quantization-Aware Training

Train the model taking quantization into consideration


- A full precision copy of the weights W is maintained throughout the training.
- The small gradients are accumulated without loss of precision.
- Once the model is trained, only the quantized weights are used for inference.

“Simulated/Fake Quantization”



Linear Quantization

An affine mapping of integers to real numbers $r = S(q - Z)$

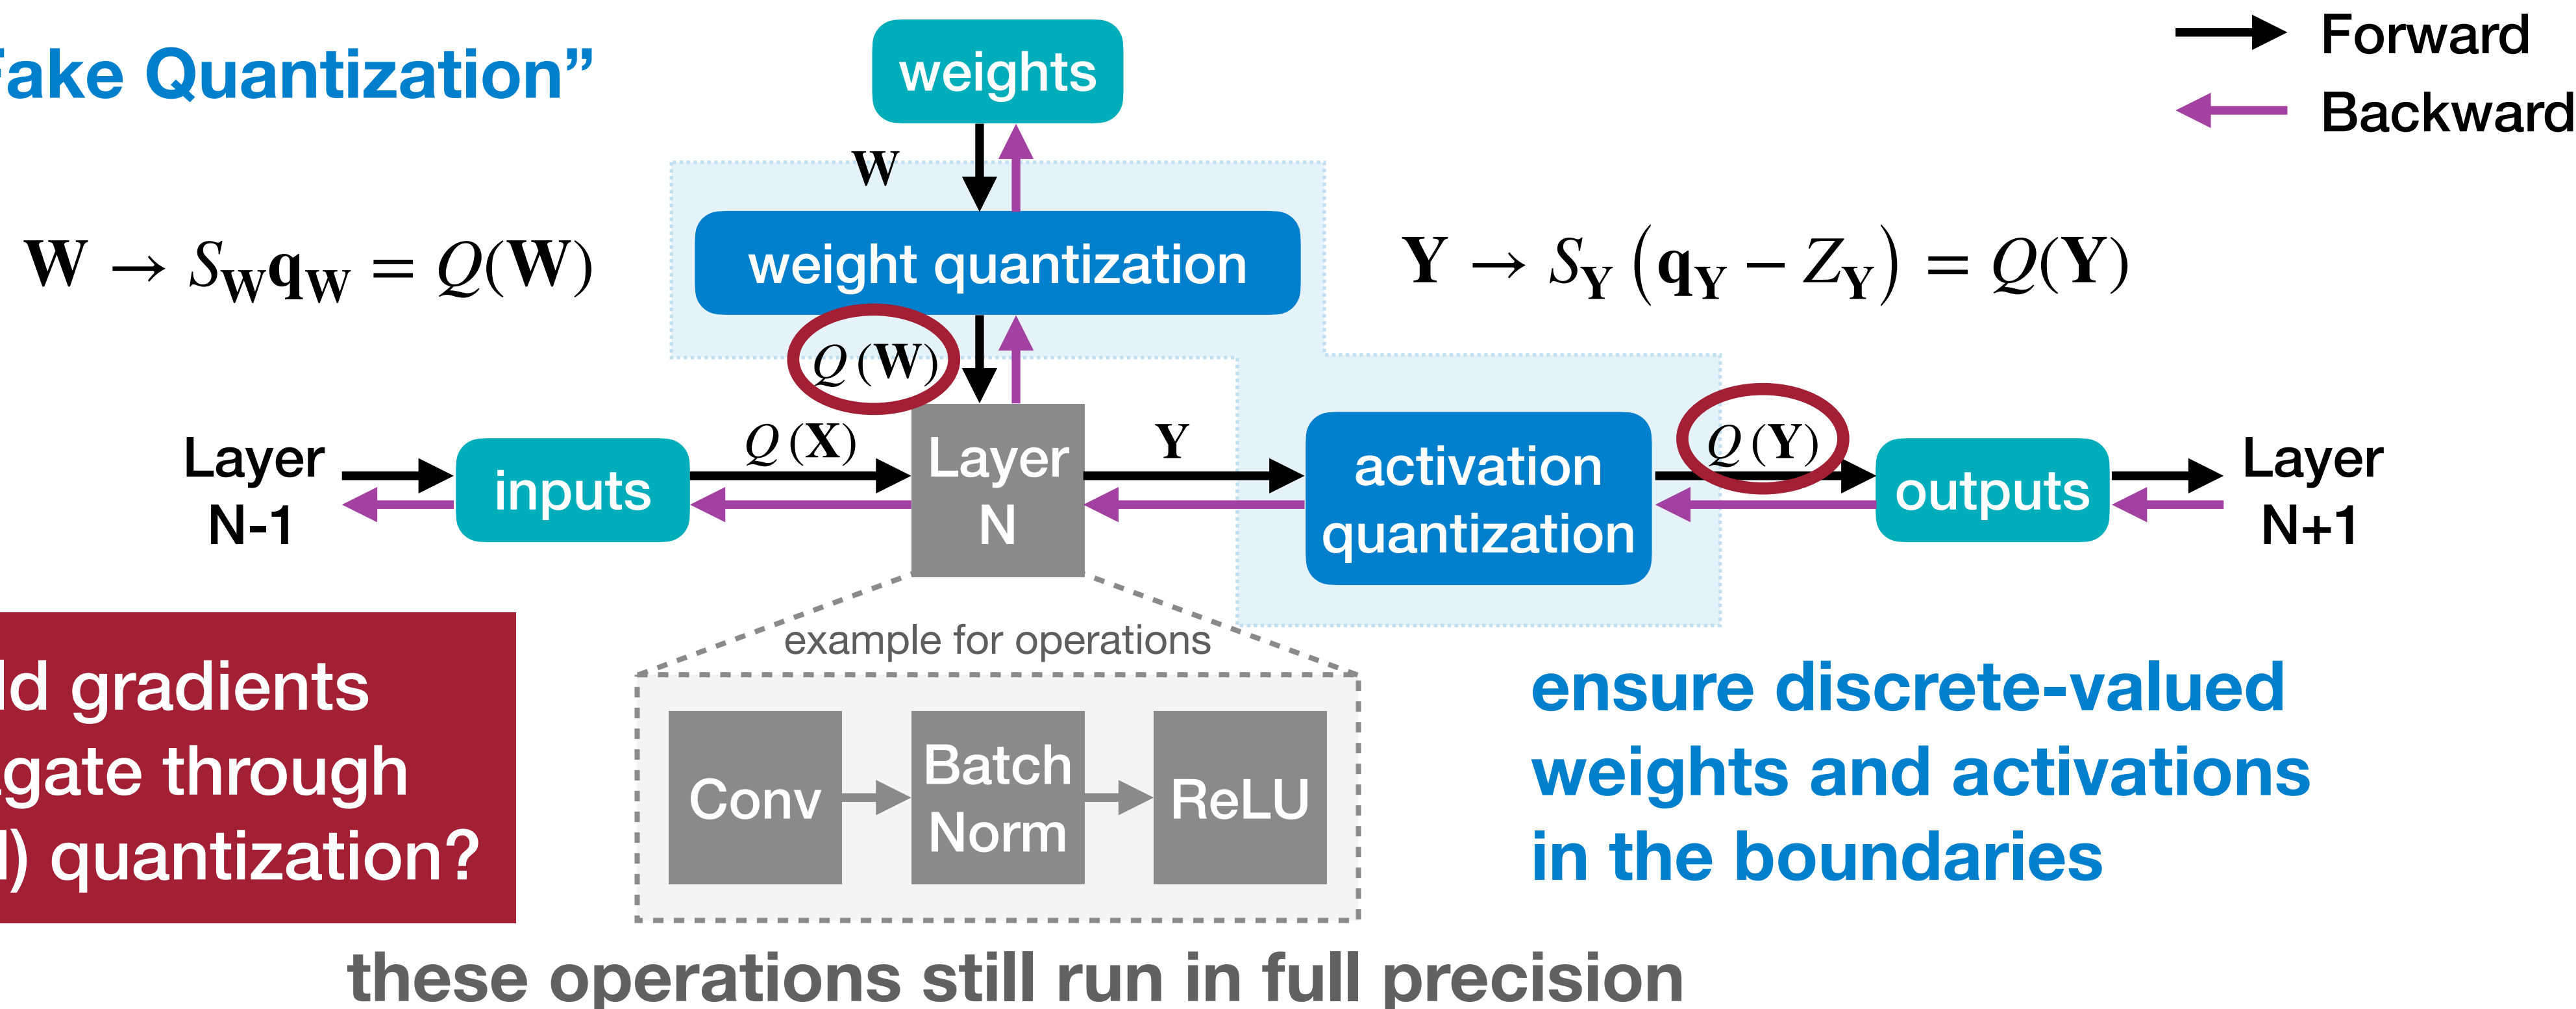


Quantization-Aware Training

Train the model taking quantization into consideration

- A full precision copy of the weights W is maintained throughout the training.
- The small gradients are accumulated without loss of precision.
- Once the model is trained, only the quantized weights are used for inference.

“Simulated/Fake Quantization”



? How should gradients back-propagate through the (simulated) quantization?

Straight-Through Estimator (STE)

- Quantization is discrete-valued, and thus the derivative is 0 almost everywhere.



$$\frac{\partial Q(W)}{\partial W} = 0$$

- The neural network will learn nothing since gradients become 0 and the weights won't get updated.

$$g_W = \frac{\partial L}{\partial W} = \frac{\partial L}{\partial Q(W)} \cdot \frac{\partial Q(W)}{\partial W} = 0$$

- Straight-Through Estimator (STE) simply passes the gradients through the quantization as if it had been the *identity* function.

$$g_W = \frac{\partial L}{\partial W} = \frac{\partial L}{\partial Q(W)}$$



Neural Networks for Machine Learning [Hinton *et al.*, Coursera Video Lecture, 2012]


Estimating or Propagating Gradients Through Stochastic Neurons for Conditional Computation [Bengio, arXiv 2013]

Quantization-Aware Training

Train the model taking quantization into consideration

- A full precision copy of the weights is maintained throughout the training.
- The small gradients are accumulated without loss of precision.
- Once the model is trained, only the quantized weights are used for inference.

“Simulated/Fake Quantization”



INT8 Linear Quantization-Aware Training

Neural Network	Floating-Point	Post-Training Quantization		Quantization-Aware Training	
		Asymmetric	Symmetric	Asymmetric	Symmetric
		Per-Tensor	Per-Channel	Per-Tensor	Per-Channel
MobileNetV1	70.9%	0.1%	59.1%	70.0%	70.7%
MobileNetV2	71.9%	0.1%	69.8%	70.9%	71.1%
NASNet-Mobile	74.9%	72.2%	72.1%	73.0%	73.0%

Quantizing Deep Convolutional Networks for Efficient Inference: A Whitepaper [Raghuraman Krishnamoorthi, arXiv 2018]

Neural Network Quantization

2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

3	0	2	1	3:	2.00
1	1	0	3	2:	1.50
0	3	1	0	1:	0.00
3	1	2	2	0:	-1.00


1	-2	0	-1
-1	-1	-2	1
-2	1	-1	-2
1	-1	0	0

(- -1) × 1.07

K-Means-based
Quantization

Linear
Quantization

Storage	Floating-Point Weights	Integer Weights; Floating-Point Codebook	Integer Weights
Computation	Floating-Point Arithmetic	Floating-Point Arithmetic	Integer Arithmetic



Neural Network Quantization

2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

3	0	2	1	3:	2.00
1	1	0	3	2:	1.50
0	3	1	0	1:	0.00
3	1	2	2	0:	-1.00

1	-2	0	-1
-1	-1	-2	1
-2	1	-1	-2
1	-1	0	0

(- -1) × 1.07

1	0	1	1
1	0	0	1
0	1	1	0
1	1	1	1

K-Means-based
Quantization

Linear
Quantization

Binary/Ternary
Quantization

Storage	Floating-Point Weights	Integer Weights; Floating-Point Codebook	Integer Weights	Binary/Ternary Weights
Computation	Floating-Point Arithmetic	Floating-Point Arithmetic	Integer Arithmetic	Bit Operations

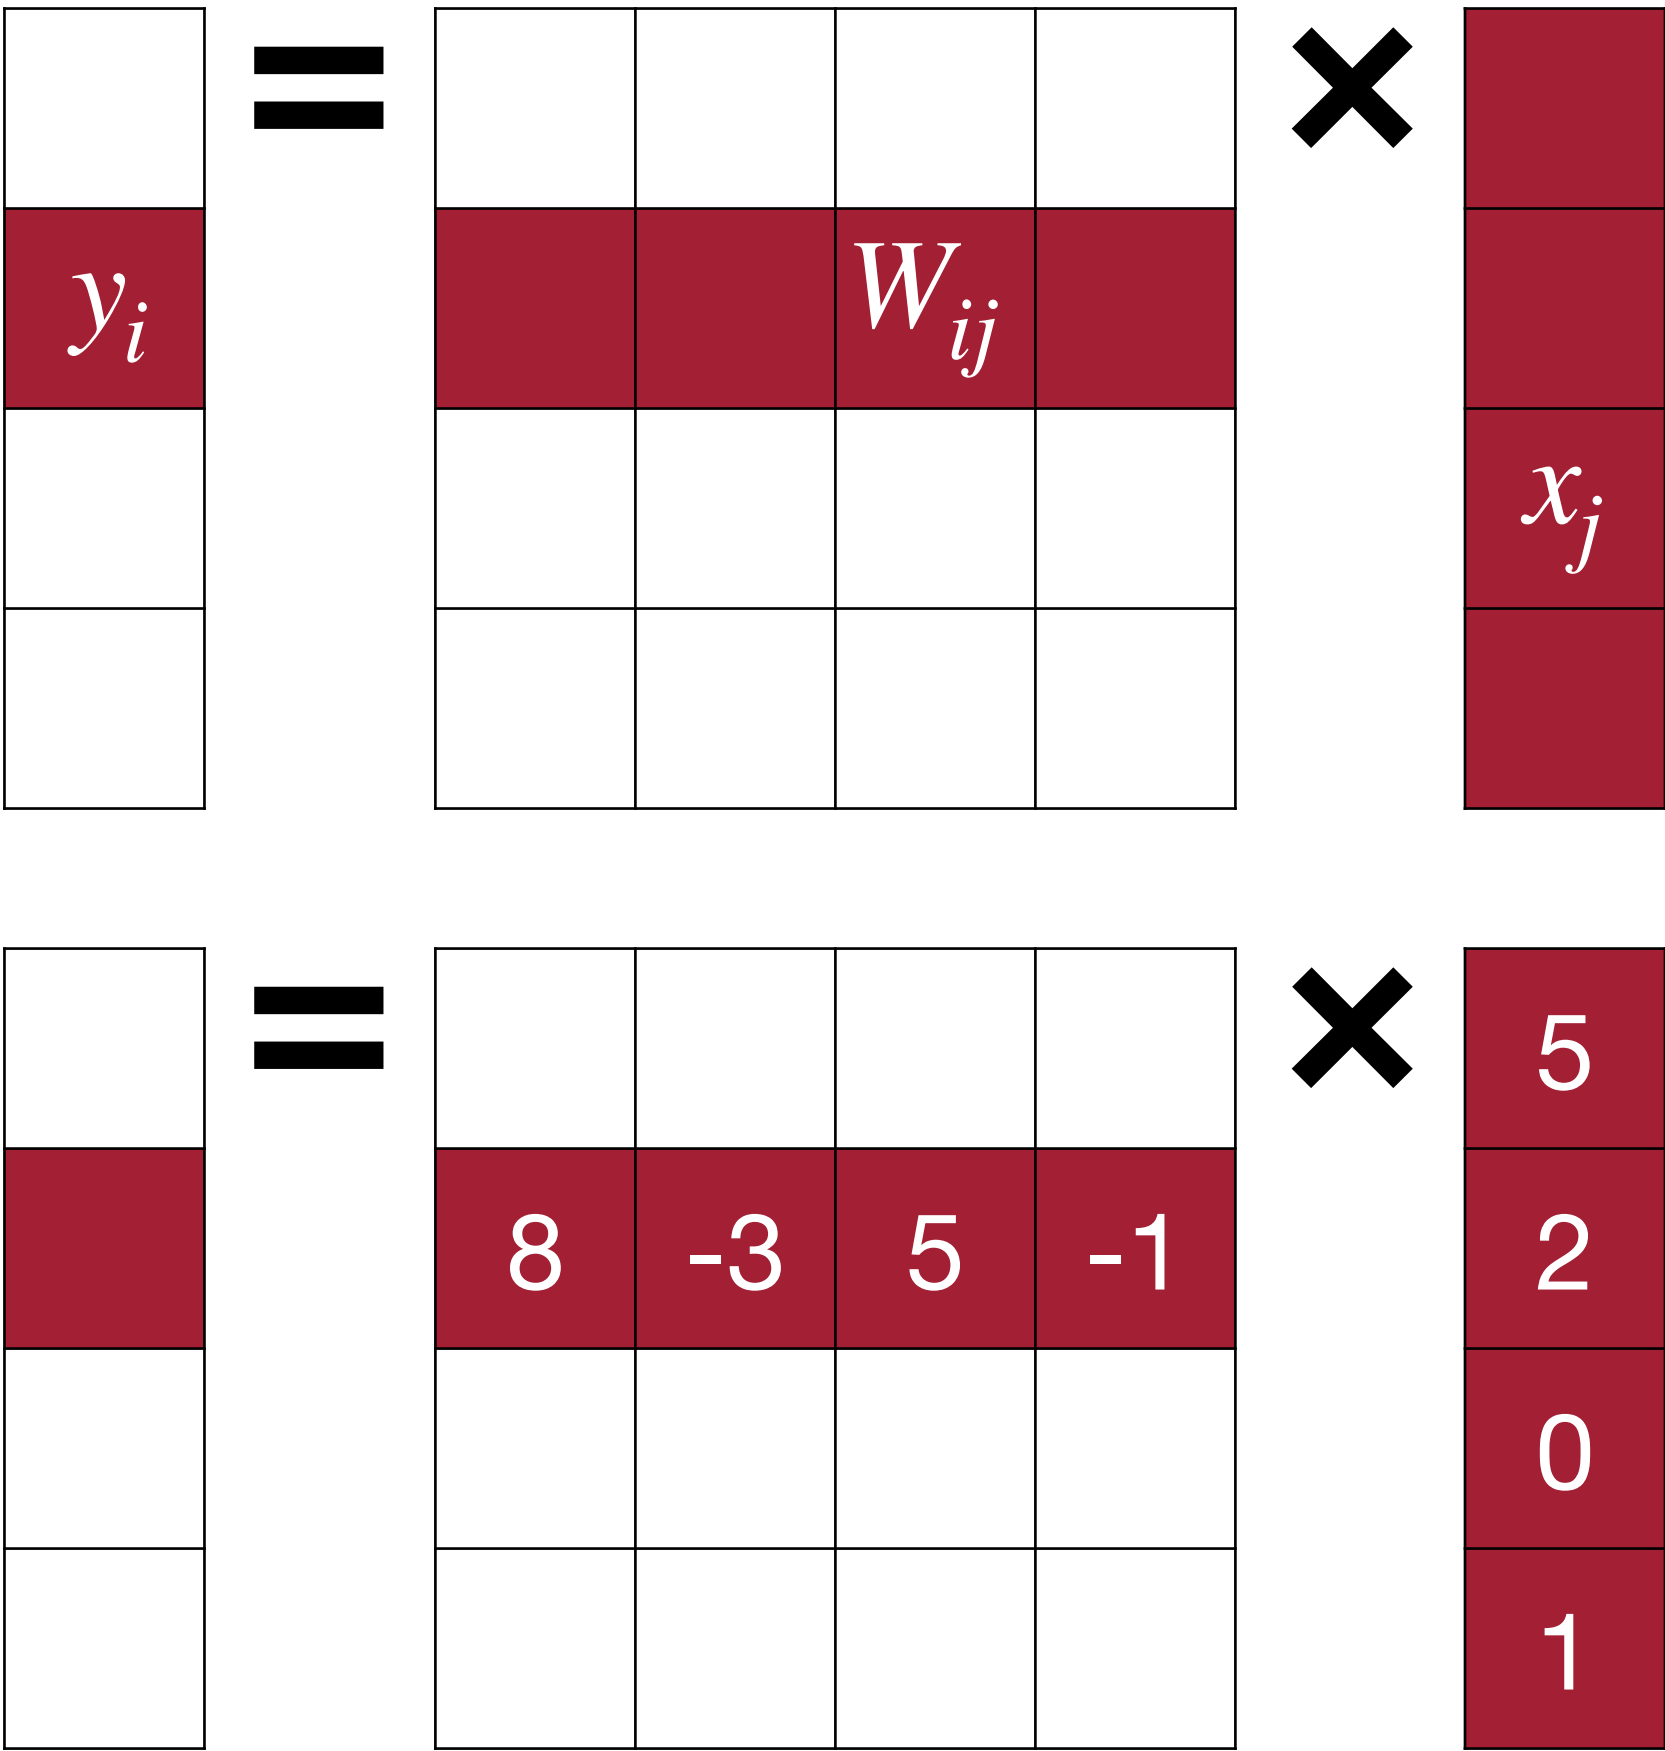
Binary/Ternary Quantization

Can we push the quantization precision to 1 bit?

Can quantization bit width go even lower?


$$y_i = \sum_j W_{ij} \cdot x_j$$
$$= 8 \times 5 + (-3) \times 2 + 5 \times 0 + (-1) \times 1$$

input	weight	operations	memory	computation
\mathbb{R}	\mathbb{R}	$+ \times$	$1 \times$	$1 \times$




If weights are quantized to +1 and -1

$$y_i = \sum_j W_{ij} \cdot x_j$$
$$= 5 - 2 + 0 - 1$$



input	weight	operations	memory	computation
\mathbb{R}	\mathbb{R}	$+ \times$	$1\times$	$1\times$
\mathbb{R}	\mathbb{B}	$+ -$	$\sim 32\times$ less	$\sim 2\times$ less



BinaryConnect: Training Deep Neural Networks with Binary Weights during Propagations [Courbariaux *et al.*, NeurIPS 2015]
XNOR-Net: ImageNet Classification using Binary Convolutional Neural Networks [Rastegari *et al.*, ECCV 2016]

Binarization

- **Deterministic Binarization**

- directly computes the bit value based on a threshold, usually 0, resulting in a sign function.

$$q = \text{sign}(r) = \begin{cases} +1, & r \geq 0 \\ -1, & r < 0 \end{cases}$$

- **Stochastic Binarization**

- use global statistics or the value of input data to determine the probability of being -1 or +1

- e.g., in Binary Connect (BC), probability is determined by hard sigmoid function $\sigma(r)$

$$q = \begin{cases} +1, & \text{with probability } p = \sigma(r) \\ -1, & \text{with probability } 1 - p \end{cases}, \quad \text{where } \sigma(r) = \min(\max(\frac{r+1}{2}, 0), 1)$$



- harder to implement as it requires the hardware to generate random bits when quantizing.

BinaryConnect: Training Deep Neural Networks with Binary Weights during Propagations [Courbariaux *et al.*, NeurIPS 2015]
BinaryNet: Training Deep Neural Networks with Weights and Activations Constrained to +1 or -1. [Courbariaux *et al.*, Arxiv 2016]

Minimizing Quantization Error in Binarization

weights
(32-bit float)

2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49



binary weights
(1-bit)

1	-1	1	1
1	-1	-1	1
-1	1	1	-1
1	1	1	1

1	-1	1	1
1	-1	-1	1
-1	1	1	-1
1	1	1	1

$$\mathbf{W}^{\mathbb{B}} = \text{sign}(\mathbf{W})$$

$$\alpha = \frac{1}{n} \|\mathbf{W}\|_1$$

$$\alpha \mathbf{W}^{\mathbb{B}}$$

$$\|\mathbf{W} - \mathbf{W}^{\mathbb{B}}\|_F^2 = 9.28$$

scale
(32-bit float)

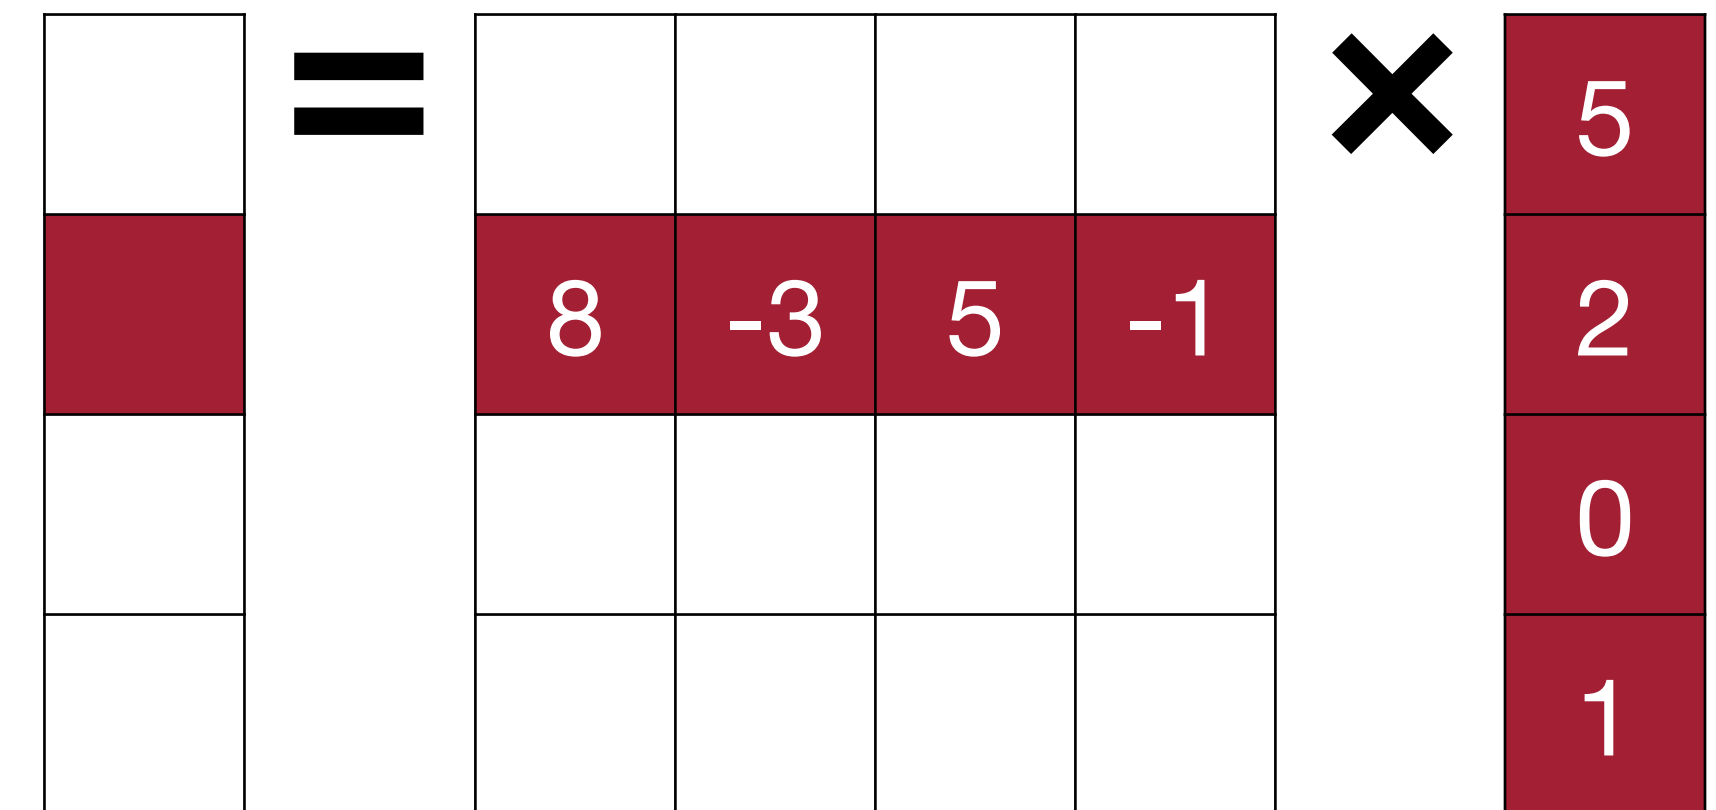
$$\times 1.05 = \frac{1}{16} \|\mathbf{W}\|_1$$

$$\|\mathbf{W} - \alpha \mathbf{W}^{\mathbb{B}}\|_F^2 = 9.24$$

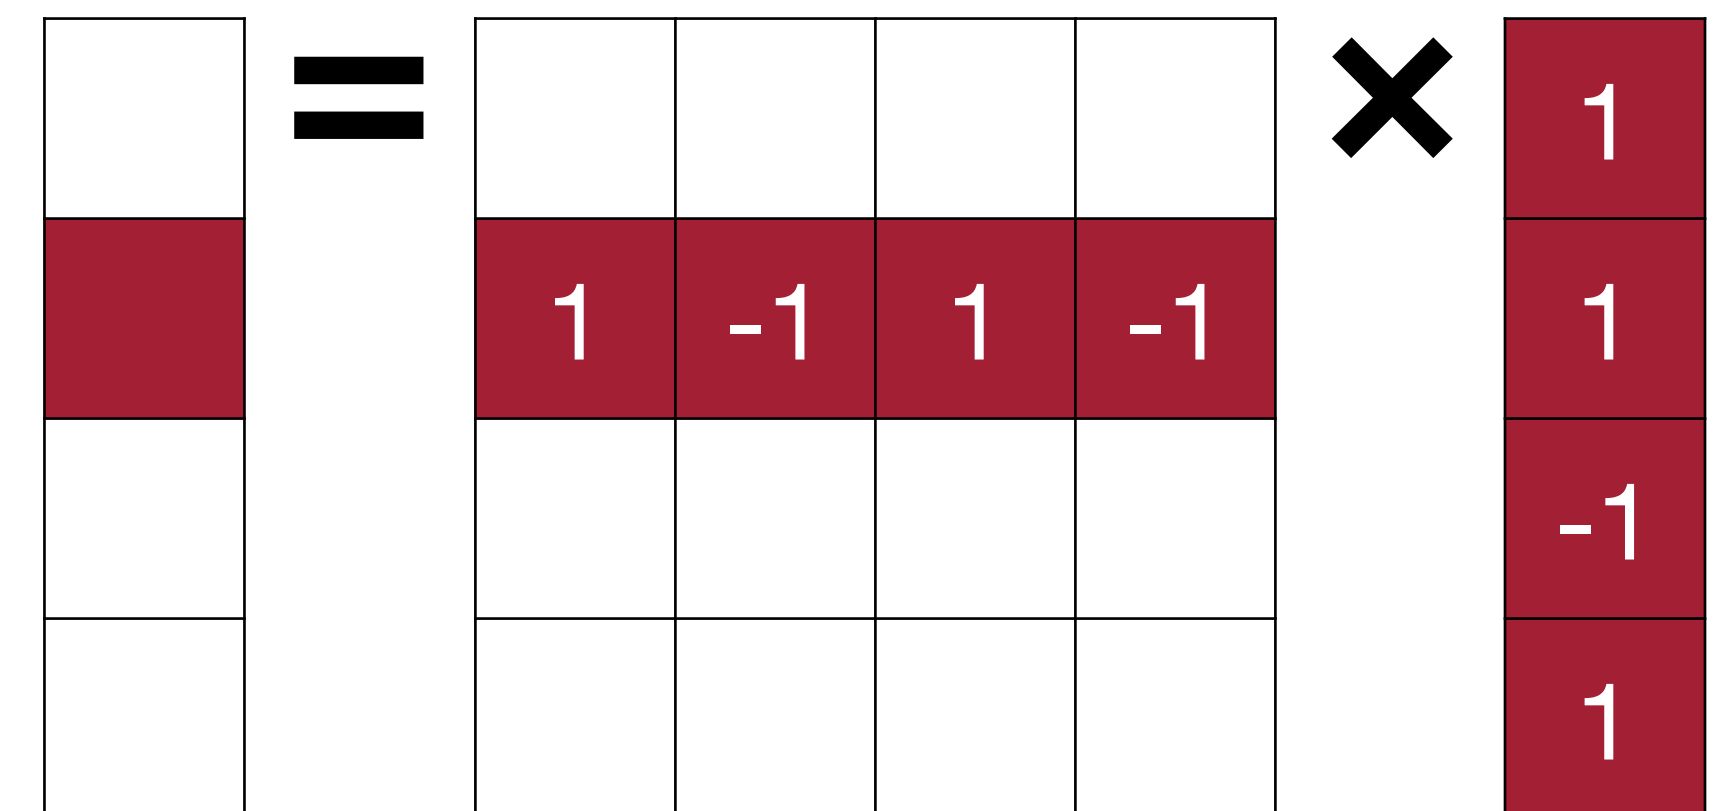
AlexNet-based Network	ImageNet Top-1 Accuracy Delta
BinaryConnect	-21.2%
Binary Weight Network (BWN)	0.2%

If both activations and weights are binarized

$$\begin{aligned}y_i &= \sum_j W_{ij} \cdot x_j \\&= 1 \times 1 + (-1) \times 1 + 1 \times (-1) + (-1) \times 1 \\&= 1 + (-1) + (-1) + (-1) = -2\end{aligned}$$



	=	<table border="1"><tr><td></td><td></td><td></td><td></td></tr><tr><td>8</td><td>-3</td><td>5</td><td>-1</td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr></table>					8	-3	5	-1									×	<table border="1"><tr><td>5</td></tr><tr><td>2</td></tr><tr><td>0</td></tr><tr><td>1</td></tr></table>	5	2	0	1
8	-3	5	-1																					
5																								
2																								
0																								
1																								



	=	<table border="1"><tr><td></td><td></td><td></td><td></td></tr><tr><td>1</td><td>-1</td><td>1</td><td>-1</td></tr><tr><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td></tr></table>					1	-1	1	-1									×	<table border="1"><tr><td>1</td></tr><tr><td>1</td></tr><tr><td>-1</td></tr><tr><td>1</td></tr></table>	1	1	-1	1
1	-1	1	-1																					
1																								
1																								
-1																								
1																								

If both activations and weights are binarized

$$y_i = \sum_j W_{ij} \cdot x_j$$
$$= 1 \times 1 + (-1) \times 1 + 1 \times (-1) + (-1) \times 1$$
$$= 1 + (-1) + (-1) + (-1) = -2$$

W	X	Y=WX
1	1	1
1	-1	-1
-1	-1	1
-1	1	-1

b _w	b _x	XNOR(b _w , b _x)
1	1	1
1	0	0
0	0	1
0	1	0

XNOR-Net: ImageNet Classification using Binary Convolutional Neural Networks [Rastegari *et al.*, ECCV 2016]

If both activations and weights are binarized

$$y_i = \sum_j W_{ij} \cdot x_j$$

$$= 1 \times 1 + (-1) \times 1 + 1 \times (-1) + (-1) \times 1$$

$$= 1 + (-1) + (-1) + (-1) = -2$$

$$= 1 \text{ xnor } 1 + 0 \text{ xnor } 1 + 1 \text{ xnor } 0 + 0 \text{ xnor } 1$$

$$= 1 + 0 + 0 + 0 = 1$$



W	X	Y=WX
1	1	1
1	-1	-1
-1	-1	1
-1	1	-1

bw	bx	XNOR(bw, bx)
1	1	1
1	0	0
0	0	1
0	1	0

XNOR-Net: ImageNet Classification using Binary Convolutional Neural Networks [Rastegari *et al.*, ECCV 2016]

If both activations and weights are binarized

$$y_i = \sum_j W_{ij} \cdot x_j$$

$$= 1 \times 1 + (-1) \times 1 + 1 \times (-1) + (-1) \times 1$$

$$= 1 + (-1) + (-1) + (-1) = -2$$

$$y_i = -n + 2 \cdot \sum_j W_{ij} \text{ xnor } x_j$$

$$= 1 \text{ xnor } 1 + 0 \text{ xnor } 1 + 1 \text{ xnor } 0 + 0 \text{ xnor } 1$$

$$= 1 + 0 + 0 + 0 = \boxed{1 \times 2 + -4} = -2$$

Assuming $\overset{\uparrow +2}{-1} \quad -1 \quad -1 \quad -1 \rightarrow$

W	X	Y=WX
1	1	1
1	-1	-1
-1	-1	1
-1	1	-1

bw	bx	XNOR(bw, bx)
1	1	1
1	0	0
0	0	1
0	1	0

XNOR-Net: ImageNet Classification using Binary Convolutional Neural Networks [Rastegari et al., ECCV 2016]

If both activations and weights are binarized

$$y_i = -n + 2 \cdot \sum_j W_{ij} \text{ xnor } x_j \quad \rightarrow \quad y_i = -n + \text{popcount}(W_i \text{ xnor } x) \ll 1$$
$$= -4 + 2 \times (1 \text{ xnor } 1 + 0 \text{ xnor } 1 + 1 \text{ xnor } 0 + 0 \text{ xnor } 1)$$
$$= -4 + 2 \times (1 + 0 + 0 + 0) = -2$$

→ **popcount**: return the number of 1

W	X	Y=WX
1	1	1
1	-1	-1
-1	-1	1
-1	1	-1

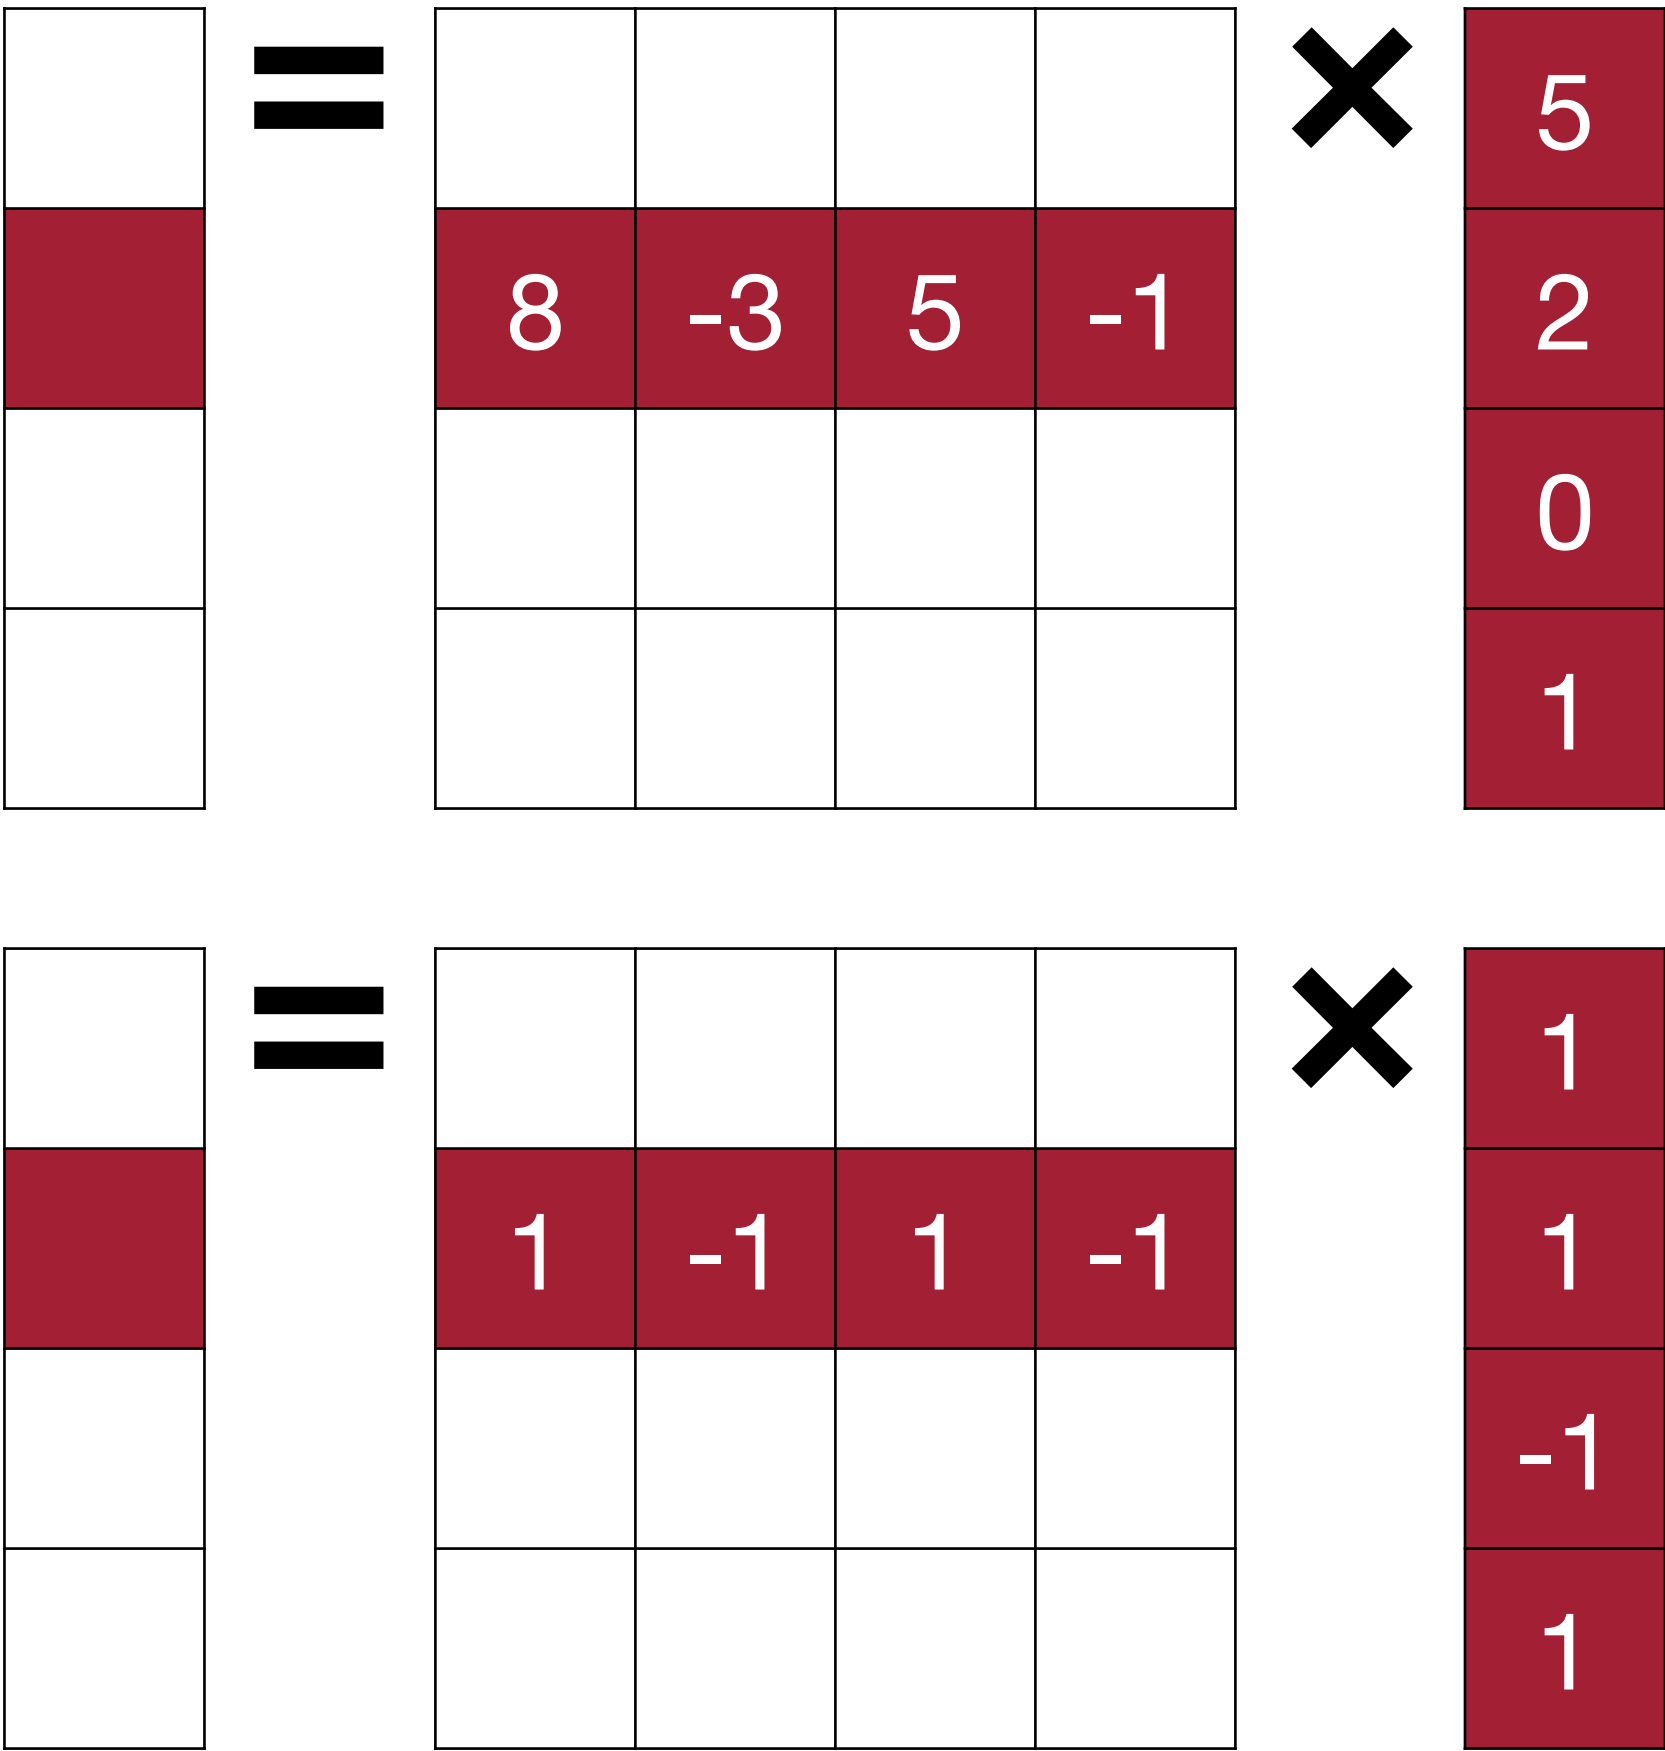
bw	bx	XNOR(bw, bx)
1	1	1
1	0	0
0	0	1
0	1	0

XNOR-Net: ImageNet Classification using Binary Convolutional Neural Networks [Rastegari *et al.*, ECCV 2016]

If both activations and weights are binarized

$$\begin{aligned} y_i &= -n + \text{popcount}(W_i \text{ xnor } x) \ll 1 \\ &= -4 + \text{popcount}(1010 \text{ xnor } 1101) \ll 1 \\ &= -4 + \text{popcount}(1000) \ll 1 = -4 + 2 = -2 \end{aligned}$$

input	weight	operations	memory	computation
\mathbb{R}	\mathbb{R}	$+ \times$	$1\times$	$1\times$
\mathbb{R}	\mathbb{B}	$+ -$	$\sim 32\times$ less	$\sim 2\times$ less
\mathbb{B}	\mathbb{B}	xnor, popcount	$\sim 32\times$ less	$\sim 58\times$ less



XNOR-Net: ImageNet Classification using Binary Convolutional Neural Networks [Rastegari et al., ECCV 2016]

Accuracy Degradation of Binarization

Neural Network	Quantization	Bit-Width		ImageNet Top-1 Accuracy Delta
		W	A	
AlexNet	BWN	1	32	0.2%
	BNN	1	1	-28.7%
	XNOR-Net	1	1	-12.4%
GoogleNet	BWN	1	32	-5.80%
	BNN	1	1	-24.20%
ResNet-18	BWN	1	32	-8.5%
	XNOR-Net	1	1	-18.1%

- * BWN: Binary Weight Network with scale for weight binarization
- * BNN: Binarized Neural Network without scale factors
- * XNOR-Net: scale factors for both activation and weight binarization

Binarized Neural Networks: Training Deep Neural Networks with Weights and Activations Constrained to +1 or -1. [Courbariaux *et al.*, Arxiv 2016]
XNOR-Net: ImageNet Classification using Binary Convolutional Neural Networks [Rastegari *et al.*, ECCV 2016]

Ternary Weight Networks (TWN)

Weights are quantized to +1, -1 and 0

$$q = \begin{cases} r_t, & r > \Delta \\ 0, & |r| \leq \Delta, \\ -r_t, & r < -\Delta \end{cases} \quad \text{where } \Delta = 0.7 \times \mathbb{E}(|r|), r_t = \mathbb{E}_{|r| > \Delta}(|r|)$$

weights \mathbf{W}
(32-bit float)

2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49



ternary weights \mathbf{W}^\top
(2-bit)

1	-1	1	0
0	0	-1	1
-1	1	0	-1
1	0	1	1

$$\Delta = 0.7 \times \frac{1}{16} \|\mathbf{W}\|_1 = 0.73$$

$$\times \quad \mathbf{1.5} = \frac{1}{11} \|\mathbf{W}_{\mathbf{W}^\top \neq 0}\|_1$$

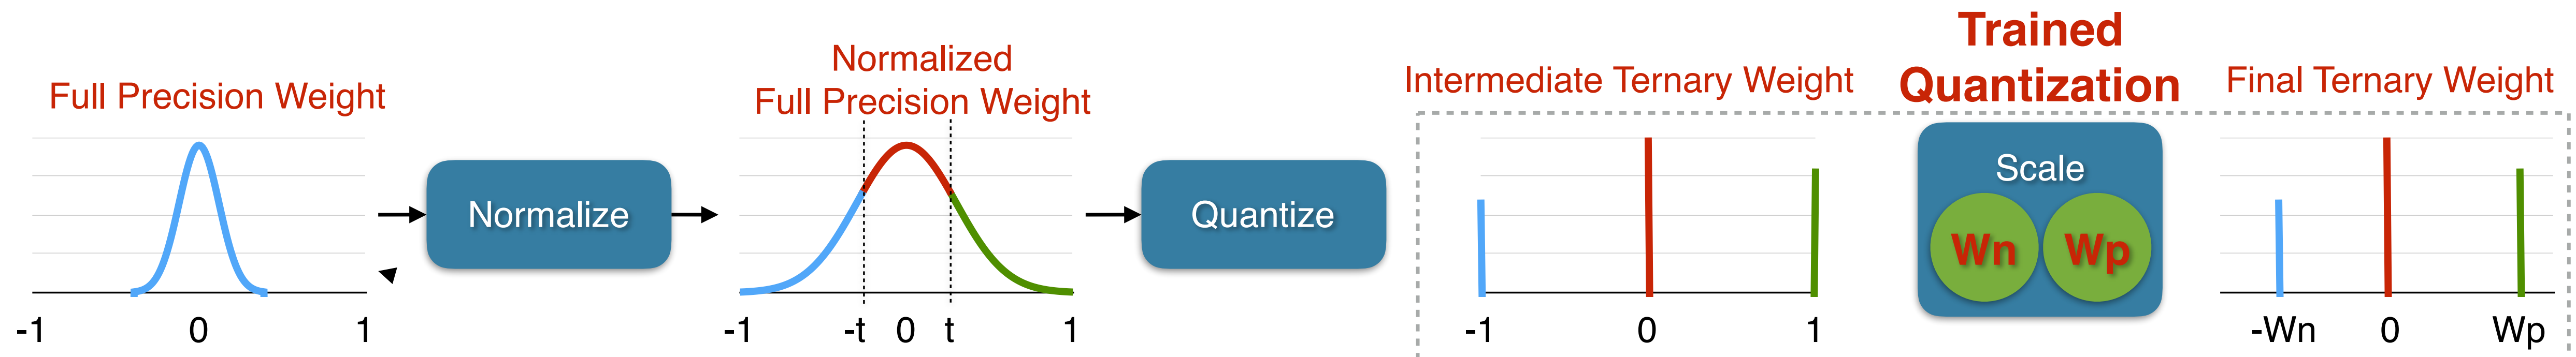
ImageNet Top-1 Accuracy	Full Precision	1 bit (BWN)	2 bit (TWN)
ResNet-18	69.6	60.8	65.3

Ternary Weight Networks [Li et al., Arxiv 2016]

Trained Ternary Quantization (TTQ)

- Instead of using fixed scale r_t , TTQ introduces two *trainable* parameters w_p and w_n to represent the positive and negative scales in the quantization.

$$q = \begin{cases} w_p, & r > \Delta \\ 0, & |r| \leq \Delta \\ -w_n, & r < -\Delta \end{cases}$$



ImageNet Top-1 Accuracy	Full Precision	1 bit (BWN)	2 bit (TWN)	TTQ
ResNet-18	69.6	60.8	65.3	66.6

Trained Ternary Quantization [Zhu *et al.*, ICLR 2017]

Neural Network Quantization

2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

3	0	2	1	3:	2.00
1	1	0	3	2:	1.50
0	3	1	0	1:	0.00
3	1	2	2	0:	-1.00

1	-2	0	-1
-1	-1	-2	1
-2	1	-1	-2
1	-1	0	0

(- -1) × 1.07

1	0	1	1
1	0	0	1
0	1	1	0
1	1	1	1

K-Means-based
Quantization

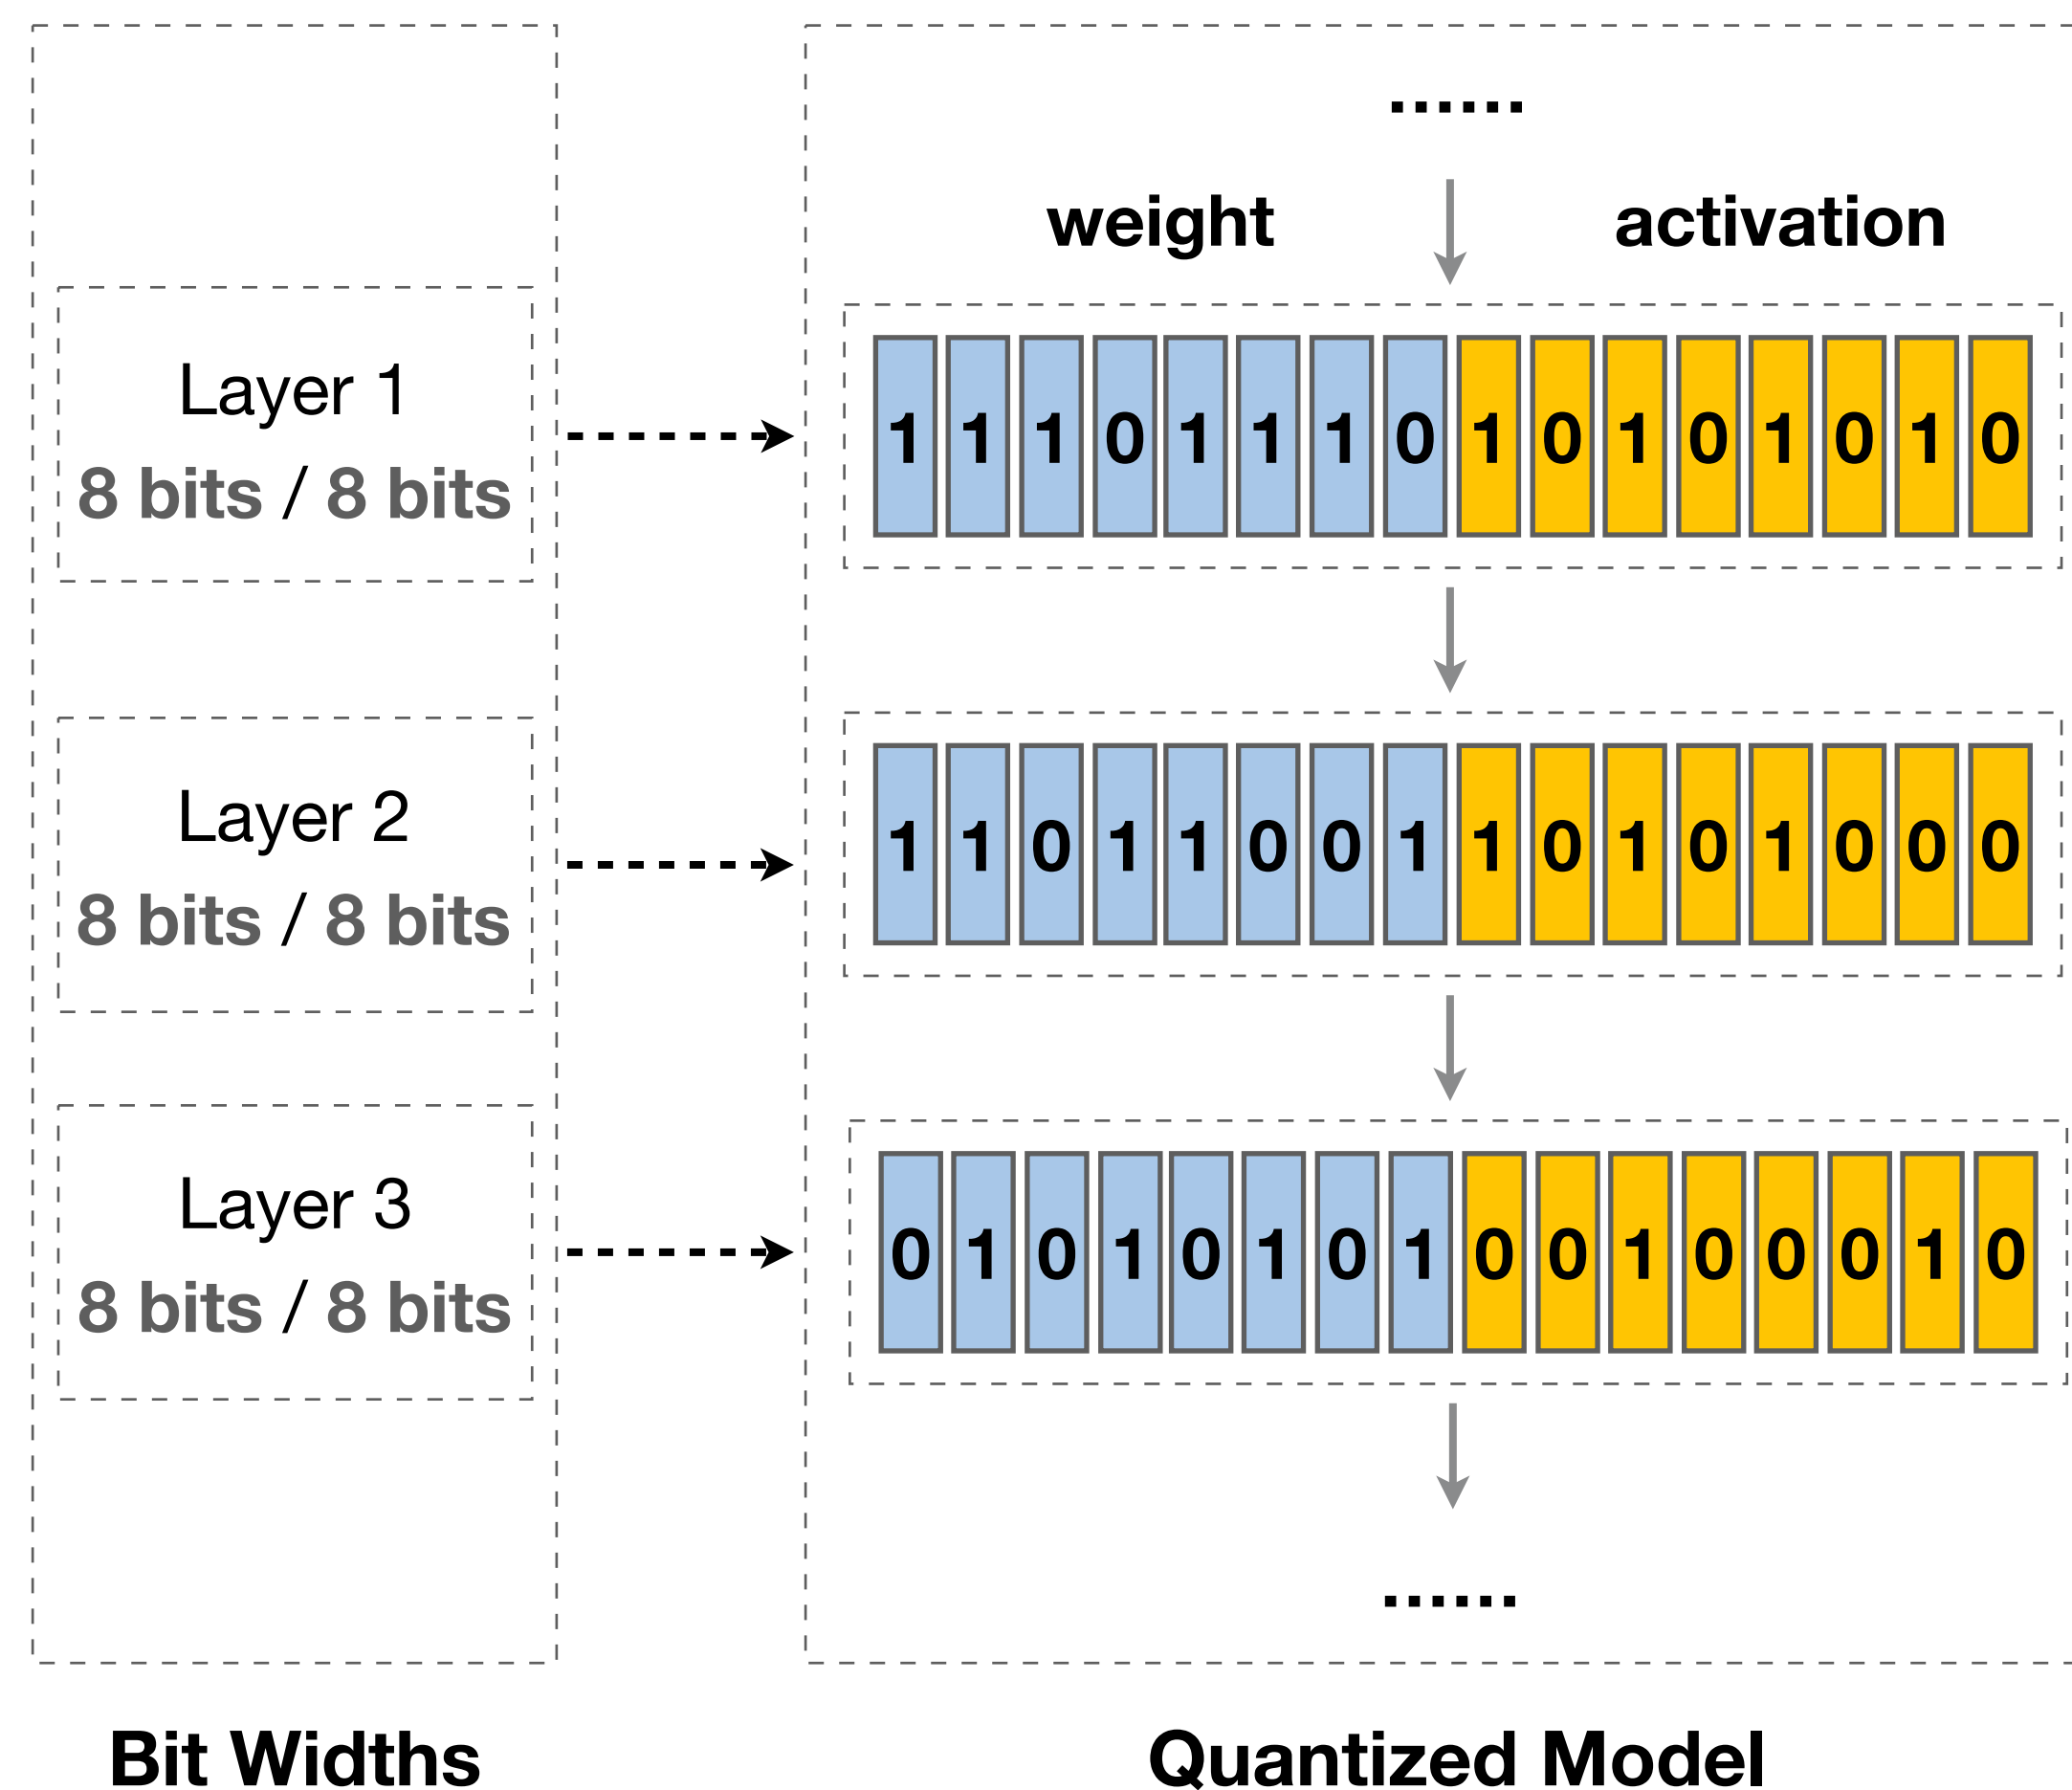
Linear
Quantization

Binary/Ternary
Quantization

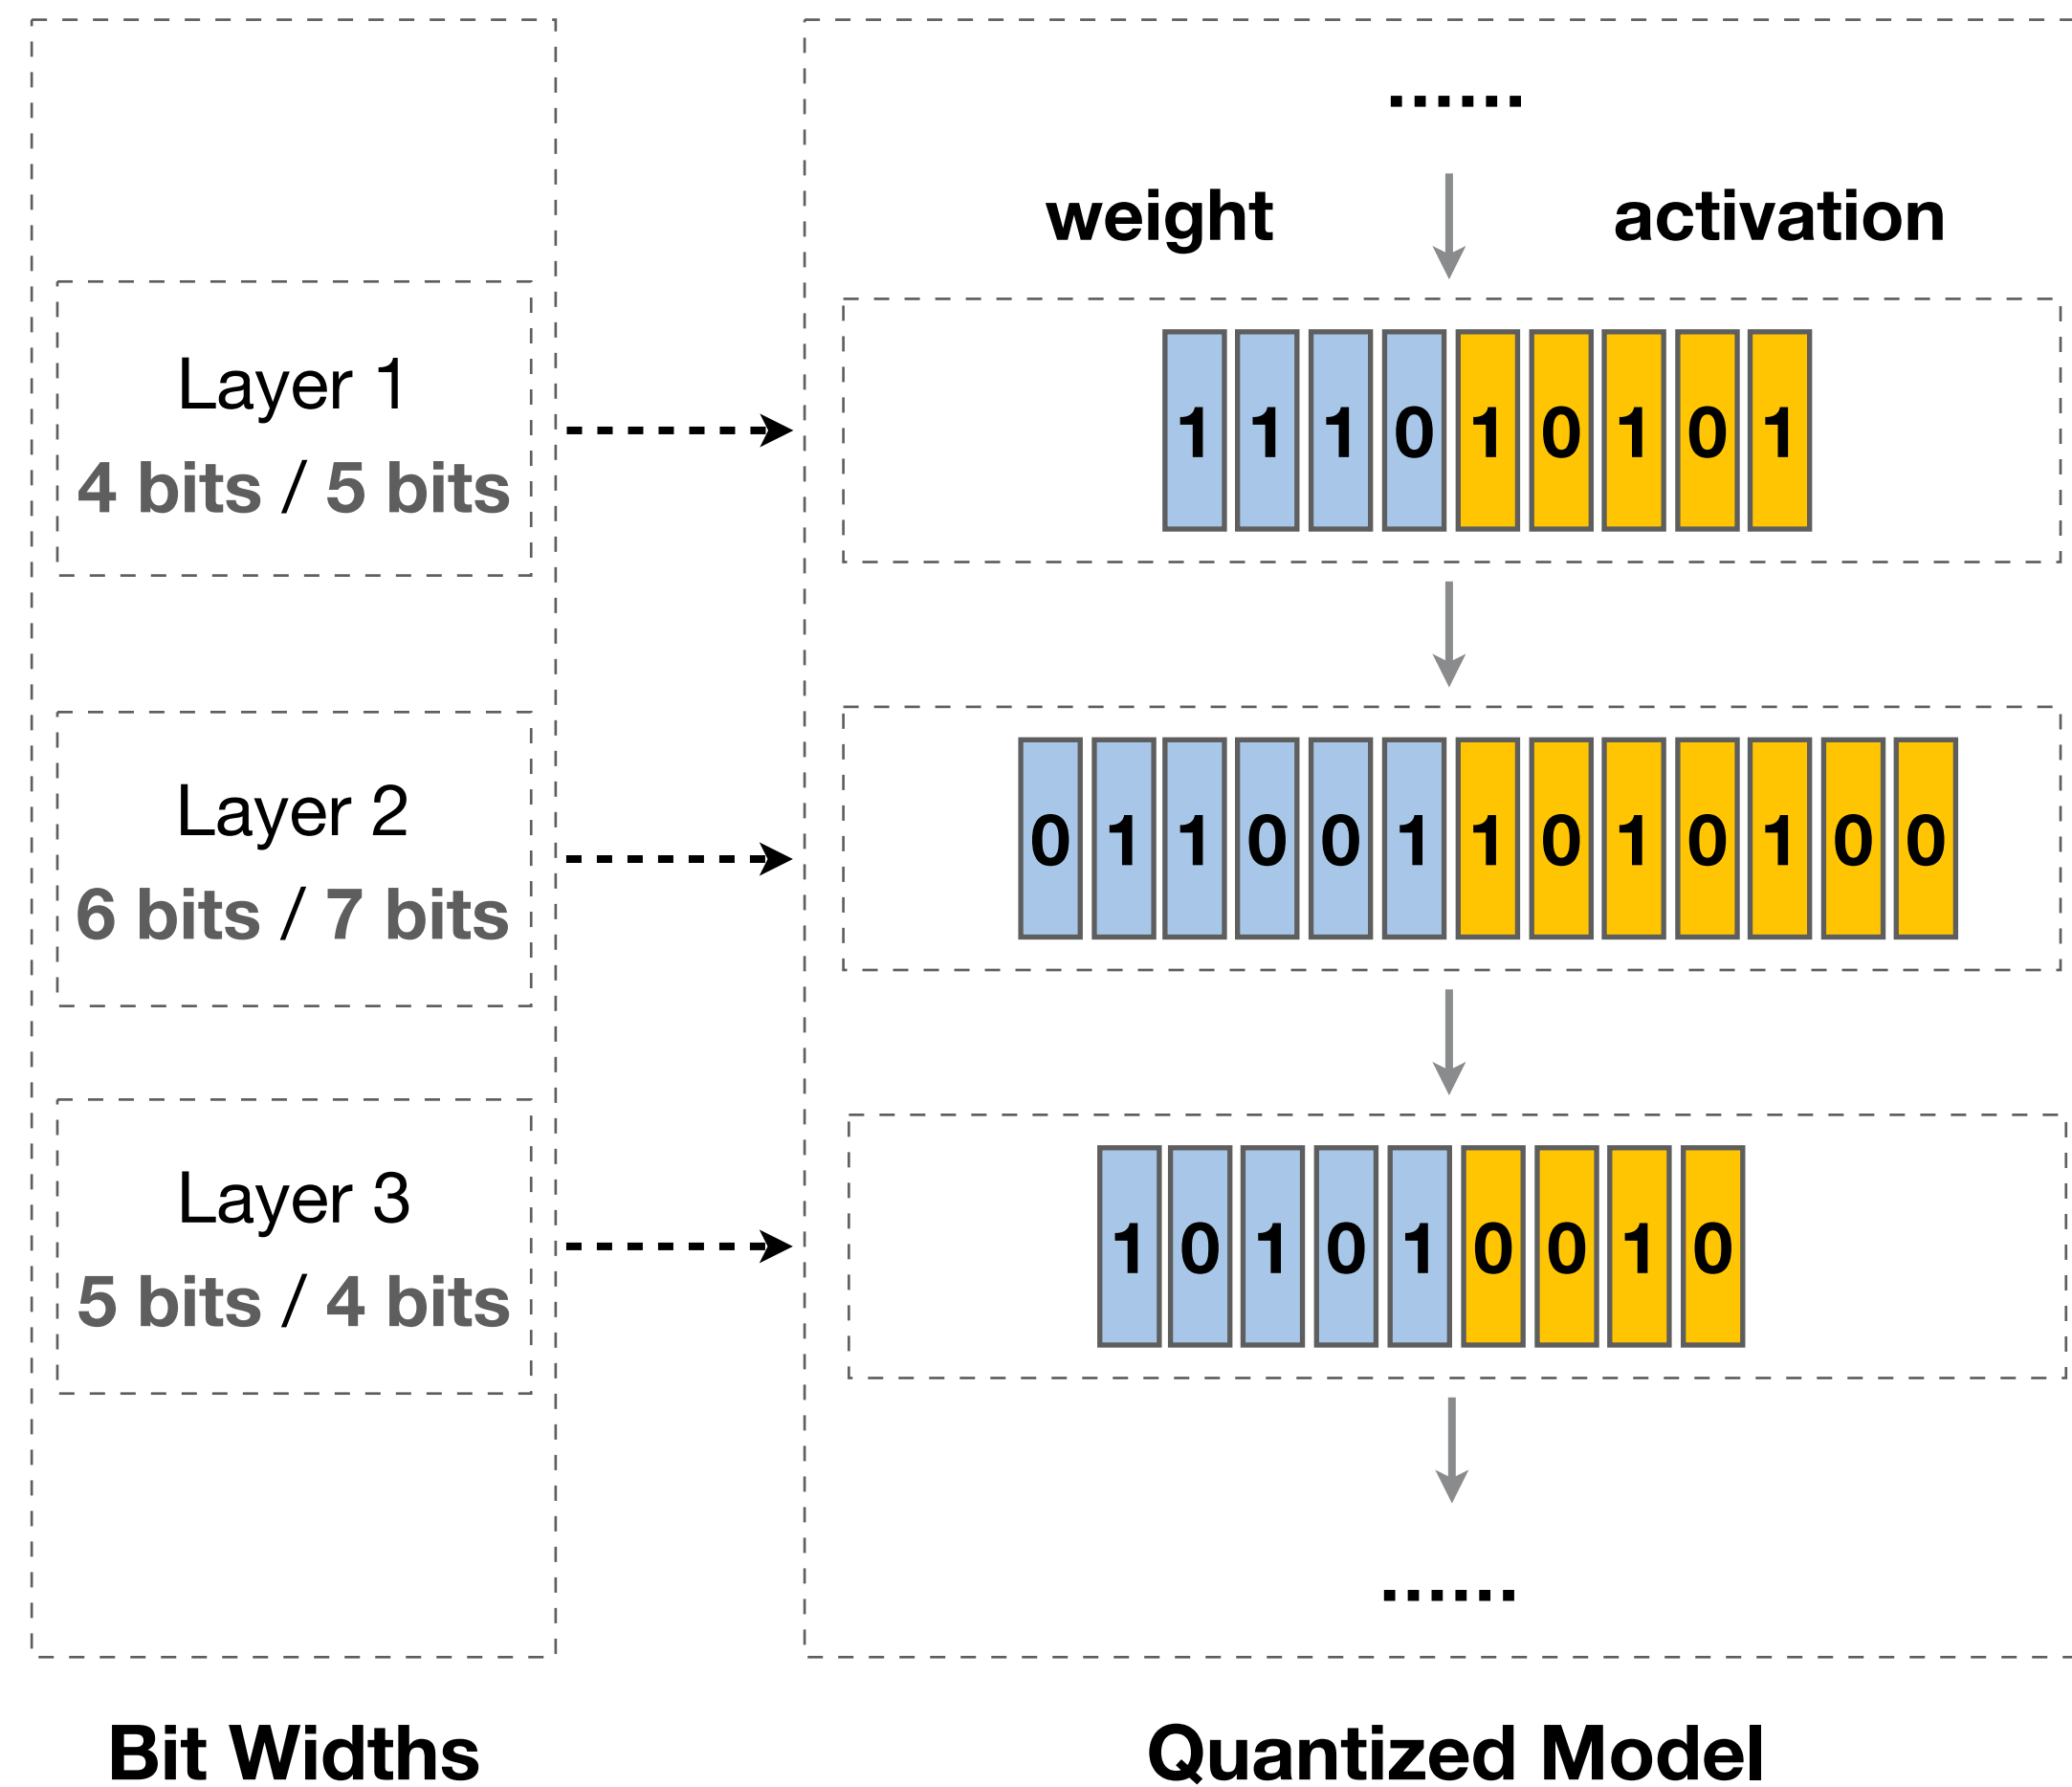
Storage	Floating-Point Weights	Integer Weights; Floating-Point Codebook	Integer Weights	Binary/Ternary Weights
Computation	Floating-Point Arithmetic	Floating-Point Arithmetic	Integer Arithmetic	Bit Operations

Mixed-Precision Quantization

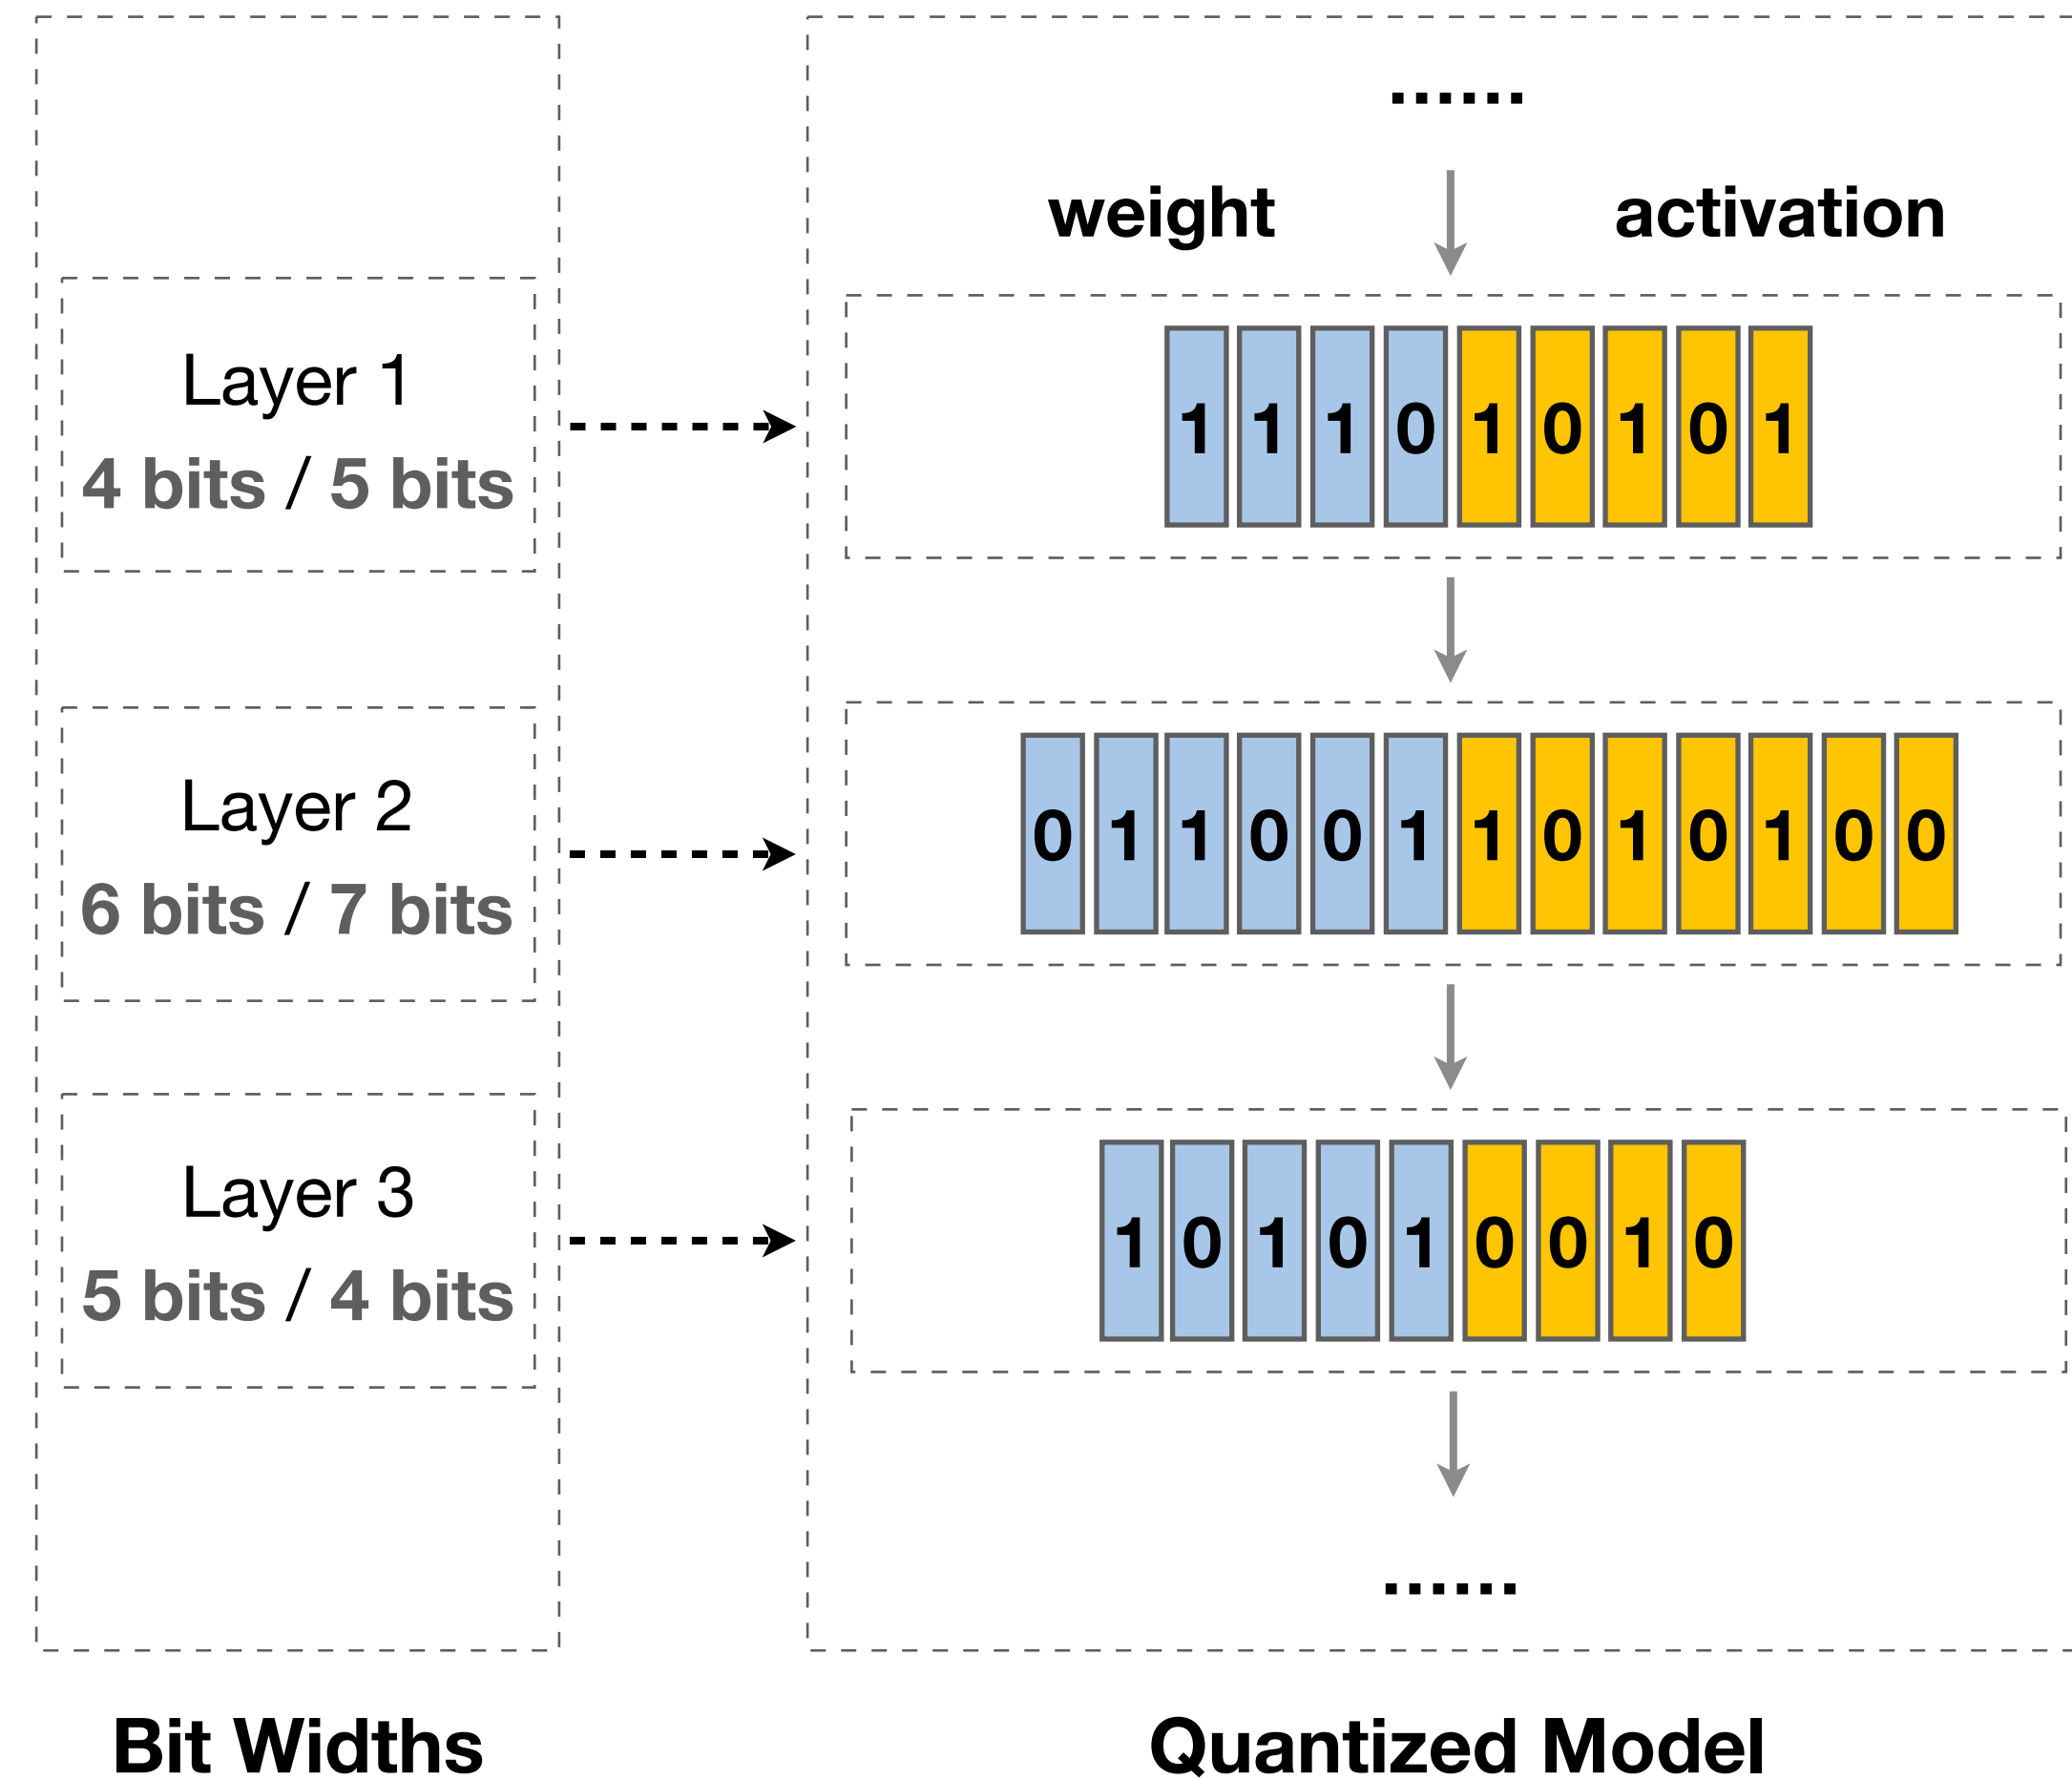
Uniform Quantization



Mixed-Precision Quantization




Challenge: Huge Design Space



Choices: $8 \times 8 = 64$


Choices: $8 \times 8 = 64$

Choices: $8 \times 8 = 64$



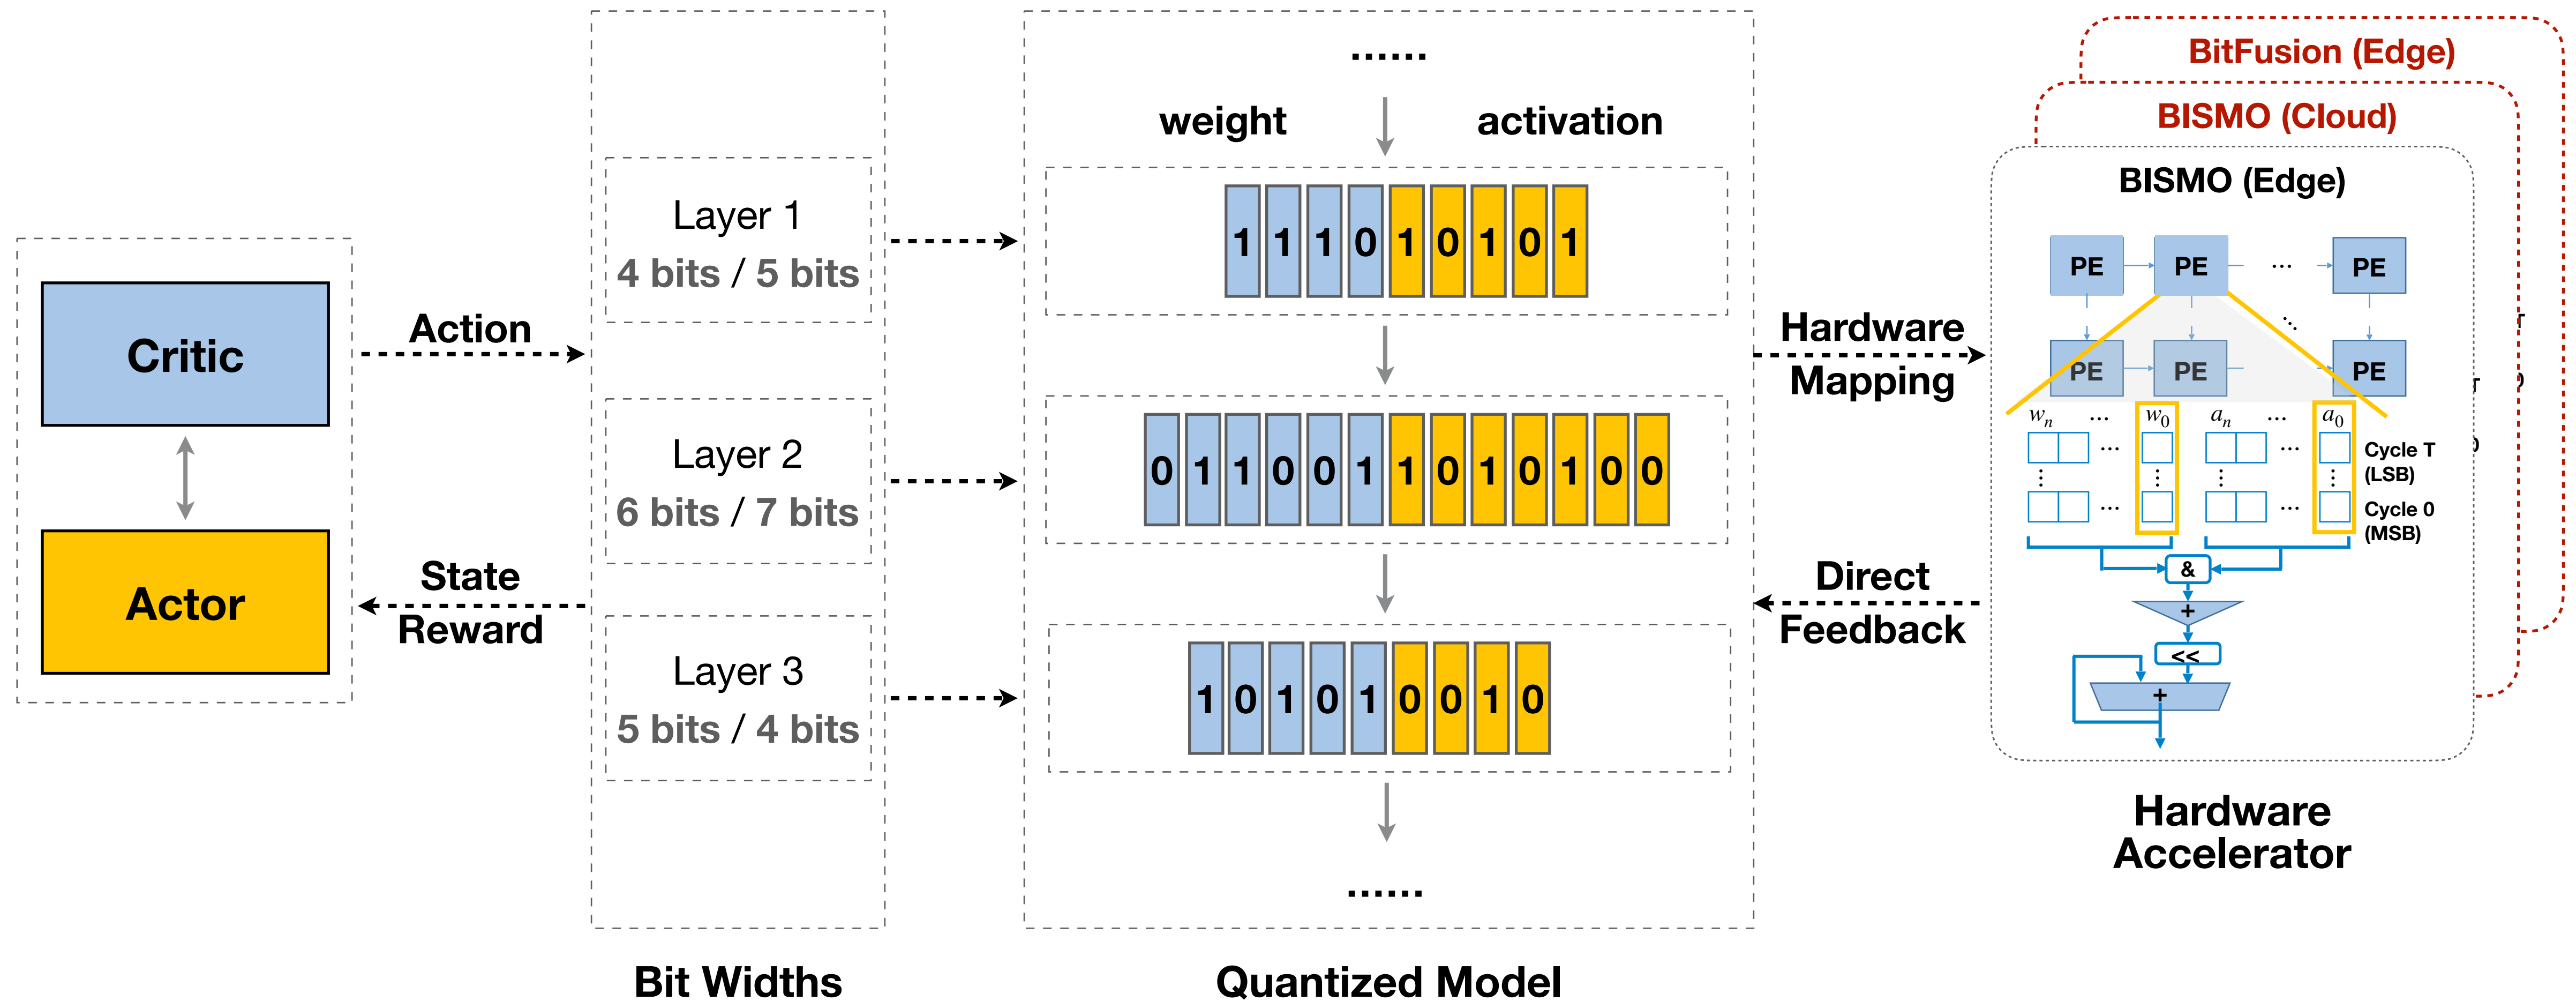
Design Space: 64^n

Solution: Design Automation




HAQ: Hardware-Aware Automated Quantization with Mixed Precision [Wang *et al.*, CVPR 2019]

Solution: Design Automation



HAQ: Hardware-Aware Automated Quantization with Mixed Precision [Wang *et al.*, CVPR 2019]

HAQ Outperforms Uniform Quantization




Mixed-Precision Quantized MobileNetV1

HAQ: Hardware-Aware Automated Quantization with Mixed Precision [Wang *et al.*, CVPR 2019]


HAQ Supports Multiple Objectives




Model Size Constrained



Latency Constrained




Energy Constrained



Mixed-Precision Quantized MobileNetV1

HAQ: Hardware-Aware Automated Quantization with Mixed Precision [Wang *et al.*, CVPR 2019]

Quantization Policy for Edge and Cloud



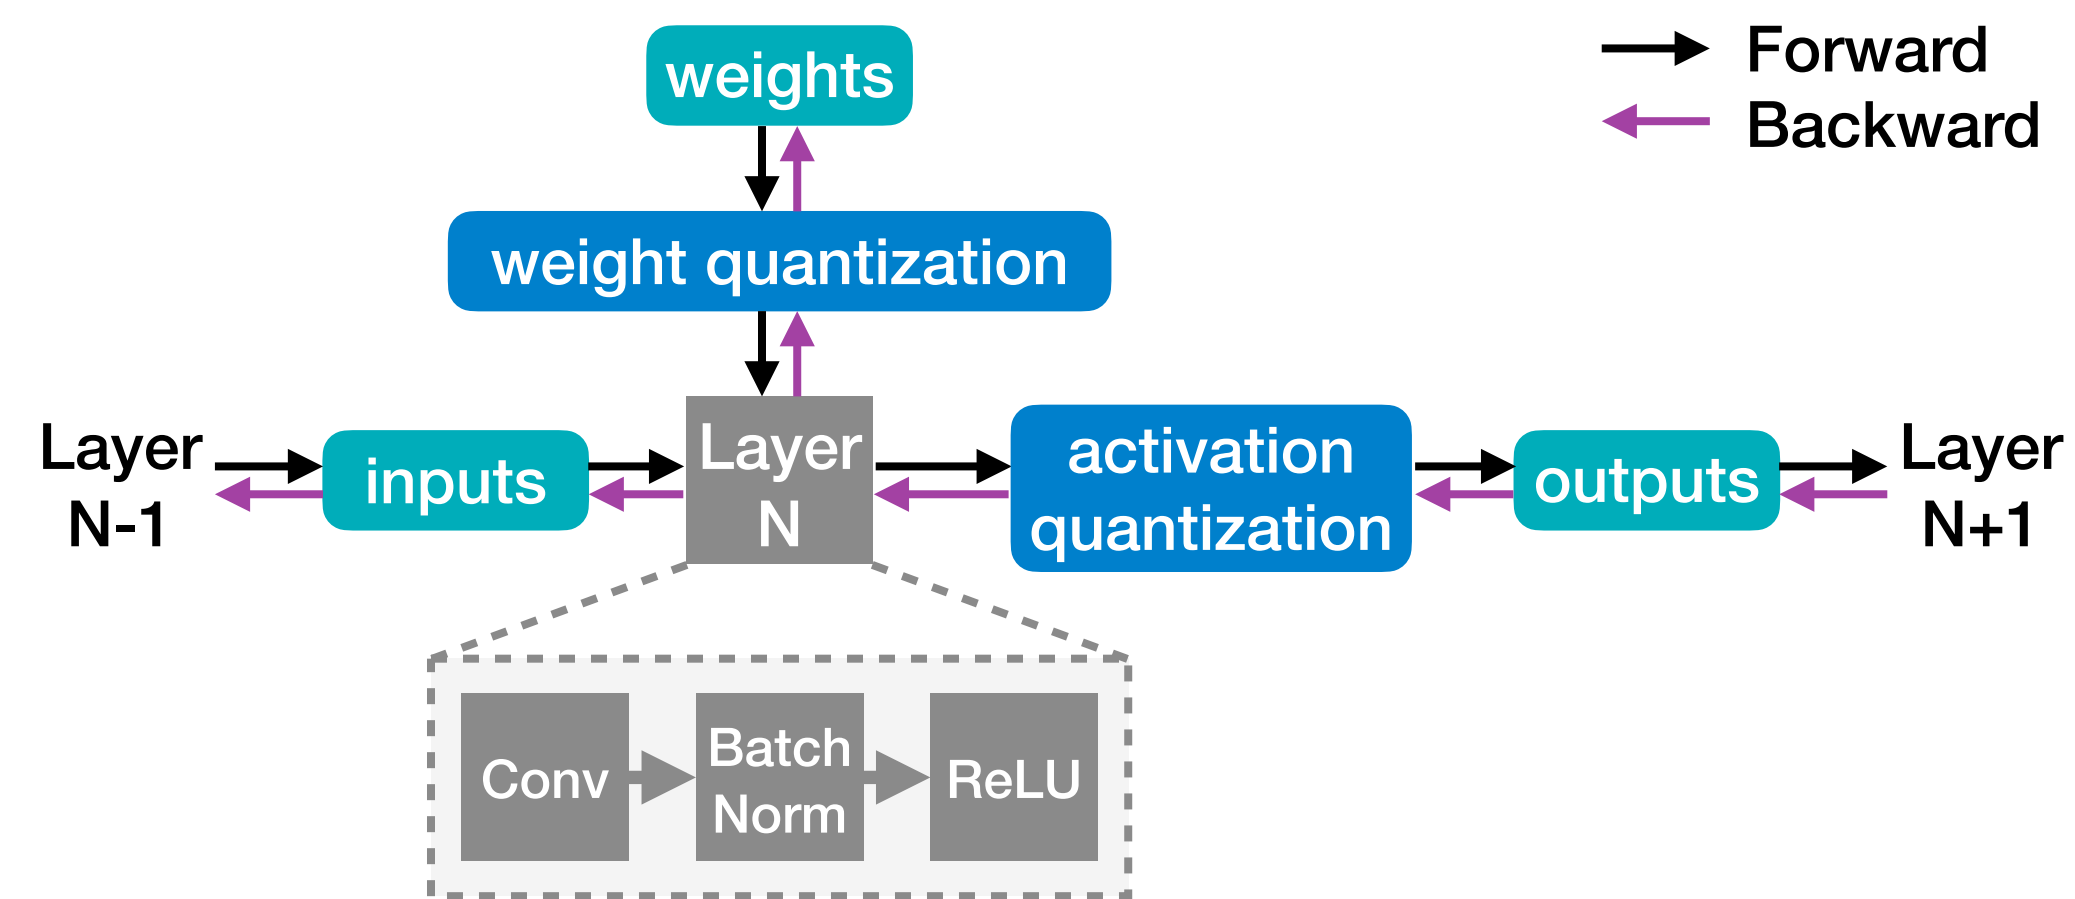

Mixed-Precision Quantized MobileNetV2

HAQ: Hardware-Aware Automated Quantization with Mixed Precision [Wang *et al.*, CVPR 2019]

Summary of Today's Lecture

In this lecture, we

1. Reviewed Linear Quantization.
2. Introduced **Post-Training Quantization (PTQ)** that quantizes an already-trained floating-point neural network model.
 - Per-tensor vs. per-channel vs. group quantization
 - How to determine dynamic range for quantization
3. Introduced **Quantization-Aware Training (QAT)** that emulates inference-time quantization during the training/fine-tuning.
 - Straight-Through Estimator (STE)
4. Introduced **binary and ternary** quantization.
5. Introduced automatic **mixed-precision** quantization.



References

1. Deep Compression [Han et al., ICLR 2016]
2. Neural Network Distiller: https://intellabs.github.io/distiller/algorithm_quantization.html
3. Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference [Jacob et al., CVPR 2018]
4. Data-Free Quantization Through Weight Equalization and Bias Correction [Markus et al., ICCV 2019]
5. Post-Training 4-Bit Quantization of Convolution Networks for Rapid-Deployment [Banner et al., NeurIPS 2019]
6. 8-bit Inference with TensorRT [Szymon Migacz, 2017]
7. Quantizing Deep Convolutional Networks for Efficient Inference: A Whitepaper [Raghuraman Krishnamoorthi, arXiv 2018]
8. Neural Networks for Machine Learning [Hinton et al., Coursera Video Lecture, 2012]
9. Estimating or Propagating Gradients Through Stochastic Neurons for Conditional Computation [Bengio, arXiv 2013]
10. Binarized Neural Networks: Training Deep Neural Networks with Weights and Activations Constrained to +1 or -1. [Courbariaux et al., Arxiv 2016]
11. DoReFa-Net: Training Low Bitwidth Convolutional Neural Networks with Low Bitwidth Gradients [Zhou et al., arXiv 2016]
12. PACT: Parameterized Clipping Activation for Quantized Neural Networks [Choi et al., arXiv 2018]
13. WRPN: Wide Reduced-Precision Networks [Mishra et al., ICLR 2018]
14. Towards Accurate Binary Convolutional Neural Network [Lin et al., NeurIPS 2017]
15. Incremental Network Quantization: Towards Lossless CNNs with Low-precision Weights [Zhou et al., ICLR 2017]
16. HAQ: Hardware-Aware Automated Quantization with Mixed Precision [Wang et al., CVPR 2019]